

Lecture 4: Threads





Lecture 4: Threads

- Overview
- Multicore Programming
- Multithreading Models
- Thread Libraries
- Implicit Threading
- Threading Issues
- Operating System Examples





Objectives

- To Introduce Notion of a **Thread**—a Fundamental Unit of CPU Utilization that Forms Basis of Multithreaded Computer Systems
- To Discuss **APIs** for **Pthreads**, Windows, and Java Thread Libraries
- To Explore Several Strategies that Provide **Implicit Threading**
- To Examine Issues Related to Multithreaded Programming
- To Cover OS Support for Threads in Windows and Linux





Motivation

- Many Parts of a Typical Application can be Run in Parallel (in **Different Threads**)
- Example 1: **Web Browser**
 - Thread A: display images
 - Thread B: retrieve data from network
- Example 2: **Word Processor**
 - Thread A: displaying graphics
 - Thread B: responding to keystrokes by user
 - Thread C: spell & grammar check in background





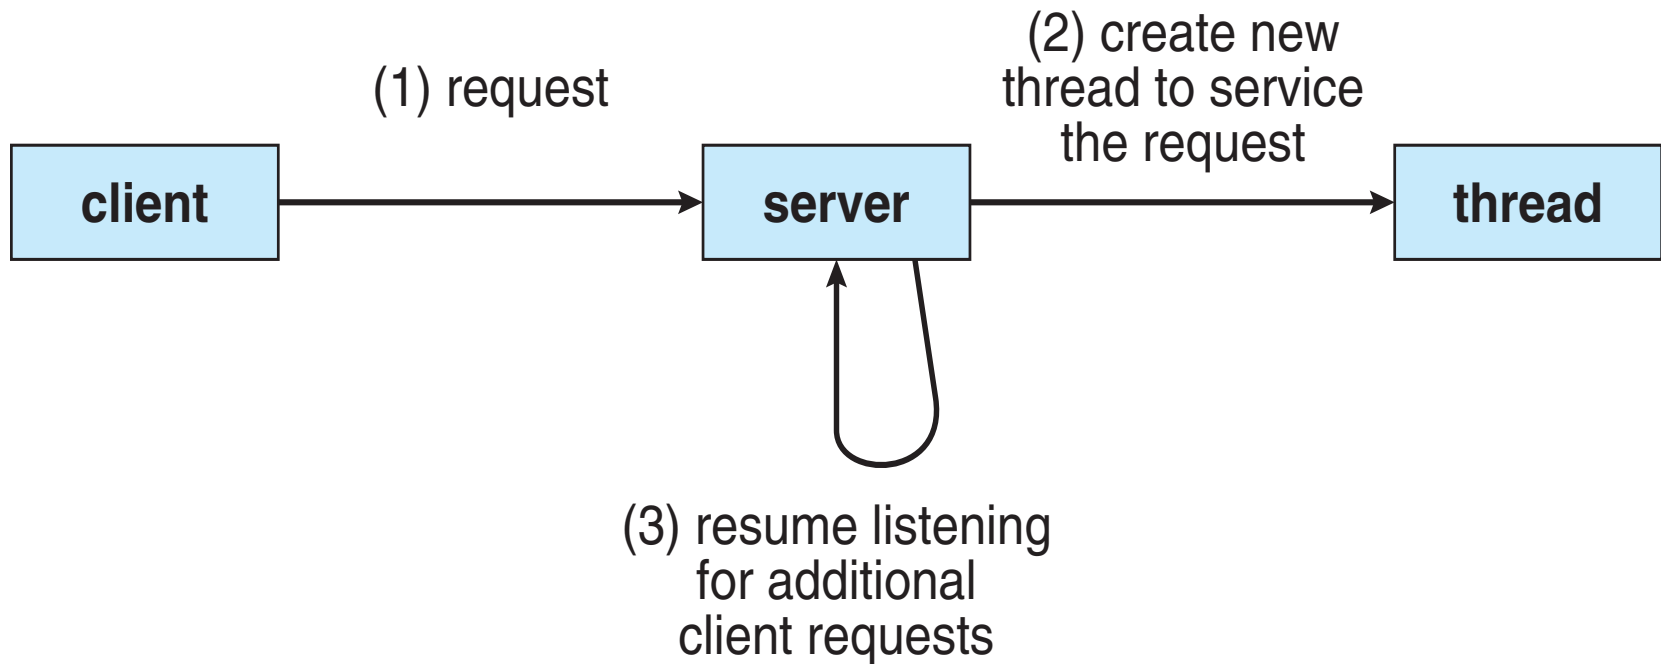
Motivation (cont.)

- Example 3: Matrix Multiplication (1000x1000)
 - Can have up to 1M parallel threads of execution
- Example 4: Remote Procedure Calls (RPCs)
- Example 5: when a process need an I/O
- Possible Solution: Creating Child Processes
 - Allows multiple flows of execution 😊
 - Time-consuming ☹️
 - Resource consumptive ☹️
- Fact:
 - Child Process will Perform Same Task as Parent Process ➔ Why to Incur such Overheads?



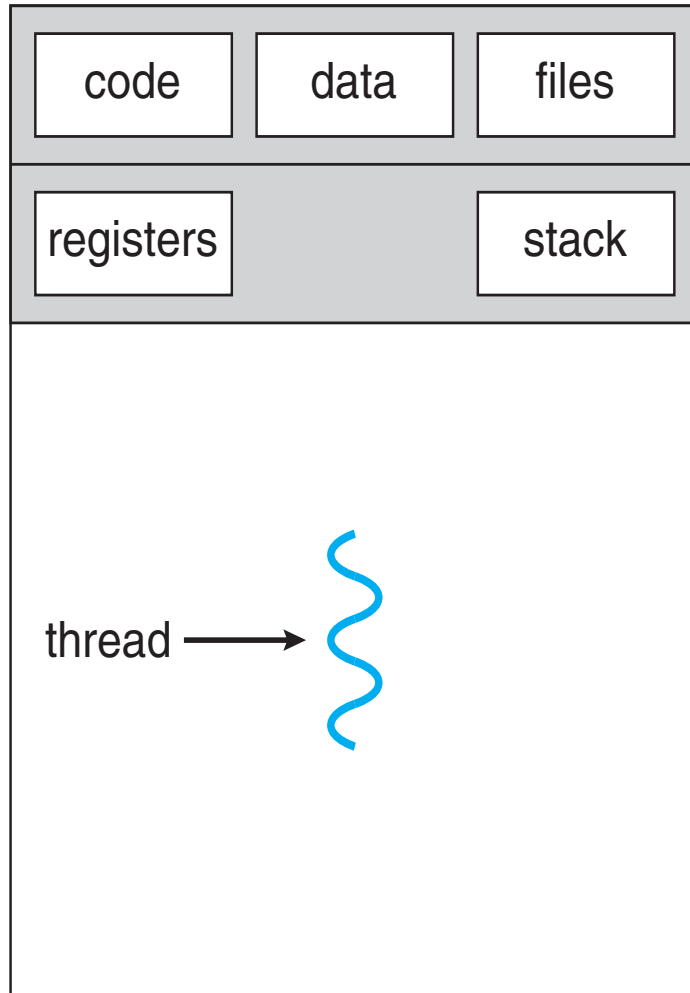


Multithreaded Server Architecture

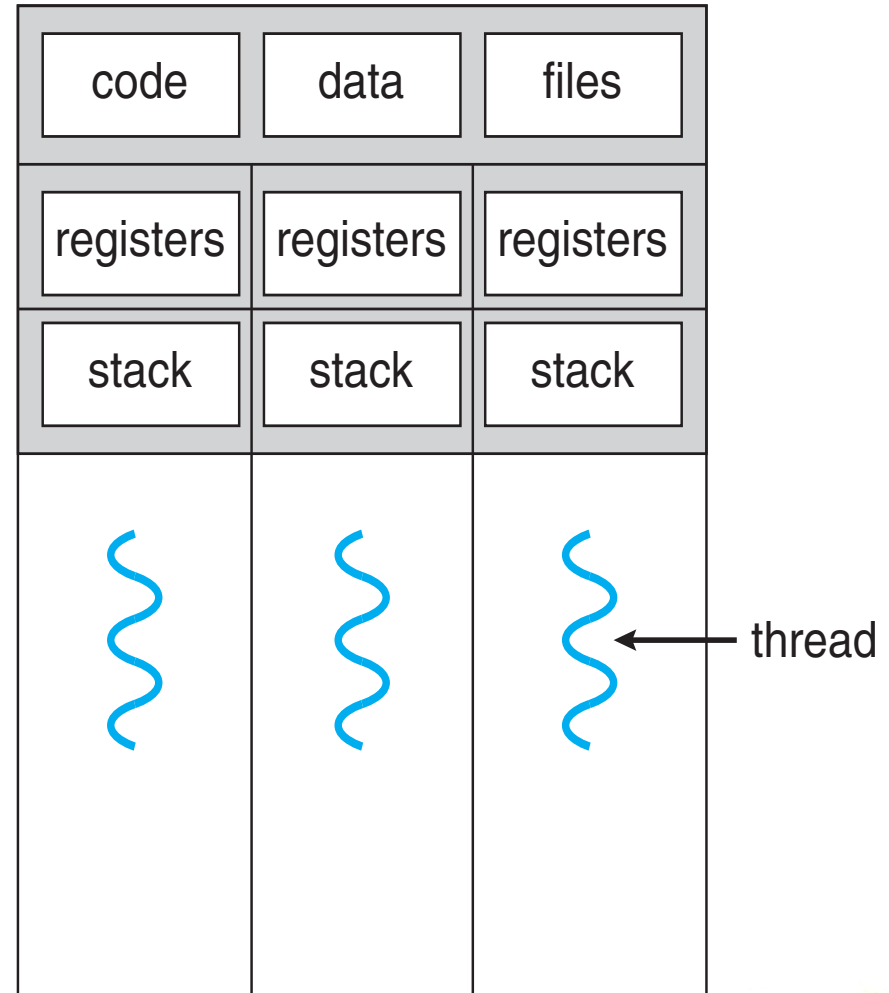




Single and Multithreaded Processes



single-threaded process



multithreaded process





Multi-Threading in OS Kernels

■ Most OSes are Multi-Threaded

- Each thread performs a specific task

■ Example


- A thread to manage I/O type A
- A thread to manage memory
- A thread for interrupt handling





Benefits of Multi-Threading

- **Responsiveness** – may allow continued execution if part of process is blocked, especially important for user interfaces
- **Resource Sharing** – threads share resources of process, easier than shared memory or message passing
- **Economy** – cheaper than process creation, thread switching lower overhead than context switching
 - Solaris: process creation 30X slower, context switching 5X slower
- **Scalability** – process can take advantage of multiprocessor architectures





Threads vs. Processes

- Processes typically **Independent**
 - Threads exist as subsets of a process
- Processes **Carry** Much more **State Info**
 - Threads of a process share process state as well as memory and other resources
- Processes have **Separate Address Spaces**
 - Threads share their address space
- Processes **Interact Only** through **IPC**
- Processes **Much Slower Context Switch Time**
- Processes **Much Slower Creation Time**





Multicore Programming

- **Multicore or Multiprocessor Systems**
Putting Pressure on Programmers
- Programming Challenges Include:
 - Dividing activities
 - Balance
 - Data splitting
 - Data dependency
 - Testing and debugging





Multicore Programming (cont.)

■ ***Parallelism***

- Implies a system can perform more than one task simultaneously

■ ***Concurrency***


- Supports more than one task making progress
- Single processor / core, scheduler providing concurrency
- Possible to have concurrency without parallelism






Concurrency vs. Parallelism

■ Concurrent Exec. on a Single-Core System:



■ Parallelism on a Multi-Core System:





Amdahl's Law

- Imagine an Application has both Serial and Parallel Components
- S is **Serial** Portion
- N Processing Cores
- **Sample** Program: 75% Parallel + 25% Serial
 - ➔ Moving from 1 to 2 Cores Results in Speedup of 1.6 times
- As N approaches infinity, speedup approaches **$1 / S$**

$$speedup \leq \frac{1}{S + \frac{(1-S)}{N}}$$





Multicore Programming (cont.)

■ Before Advent of Multi-core Systems

- CPU schedulers designed to provide **illusion of parallelisms** (not actual parallelism)
- Processes were running **concurrently** (but not in parallel)

■ As # of Threads Grows, so Does Architectural Support for Threading

- CPUs have cores as well as ***HW threads***
 - ▶ HW support such as **multiple sets of RFs and PCs**
- Consider Oracle SPARC T4 with 8 cores, and 8 hardware threads per core





User Threads and Kernel Threads

■ Kernel Threads

- Supported and managed directly by OS

■ User Threads

- Management done by user-level threads library
- Supported above kernel and managed **without** kernel support
 - ▶ **Kernel is unaware** of them unless they are mapped to kernel threads
- Context switching very fast (no interaction with kernel)

■ **Note:** Ultimately, a Relationship Must Exist between User threads and Kernel threads





Multithreading Models

■ Many-to-One (N:1)

- All application-level threads maps to a single kernel-level scheduled entity

■ One-to-One (1:1)

- Threads created by user are in 1-1 correspondence with schedulable entities in kernel (Some say, **no user-level thread**)

■ Many-to-Many (M:N)

- Maps M number of application-level threads to N number of kernel entities





Many-to-One

- Many User-Level Threads Mapped to Single Kernel Thread
- Thread management Done by thread library in user space → fast & efficient
- Entire Process Will Block if a thread Makes a **Blocking System Call**

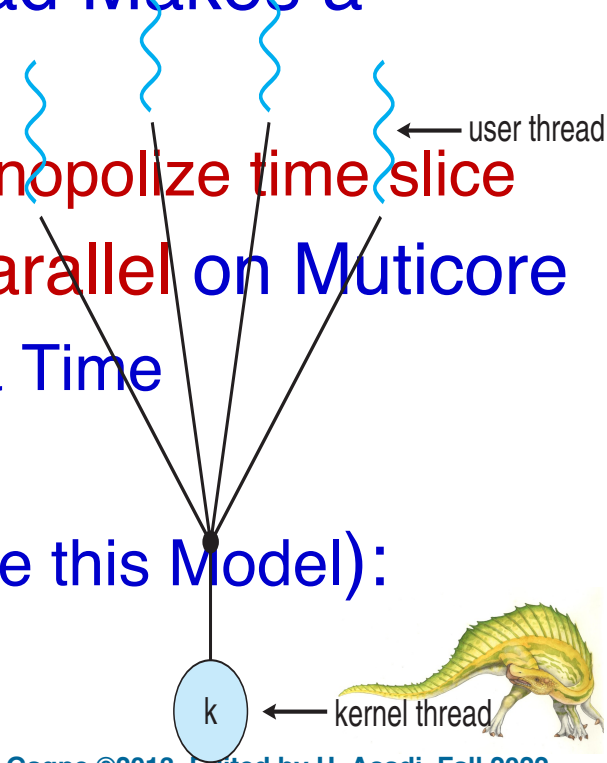
- Alternatively, a single thread can **monopolize time slice**

- Multiple Threads **may not Run in Parallel** on Multicore

- Since only one may be in Kernel at a Time
 - Does **not benefit** from **multi cores**

- Examples (Few Systems Currently use this Model):


- **Solaris Green Threads**
 - **GNU Portable Threads**





One-to-One

- Each User-Level Thread Maps to Kernel Thread
- Creating a User-Level Thread Creates a Kernel Thread
- **More Concurrency** than Many-to-One
 - Allows another thread to run when a thread makes a **blocking system call**
- **Examples**
 - Windows, Linux
 - Solaris 9 and later






One-to-One (cont.)

■ Main Drawback


- Creating a user thread requires creating corresponding kernel thread
- → Overhead of creating kernel threads can burden performance of an application
- **Solution:** most implementations of this model restrict number of threads supported by system





Many-to-Many Model

- **Multiplexes Many** User Level Threads to be Mapped to **Smaller** or Equal Number of Kernel Threads
- Allows OS to Create a Sufficient Number of Kernel Threads
- Tries to Address Shortcomings of Many-to-one and One-to-One Models
- Solaris prior to version 9
- Windows with *ThreadFiber* package






Two-level Model

■ Similar to M:M, Except that it Allows a User Thread to be **Bound** to Kernel Thread

■ Examples

- IRIX
- HP-UX
- Tru64 UNIX
- Solaris 8 and earlier





Processes, User Threads, Kernel Threads

■ Process

- Isolated with its own virtual address space
- Contains process data like file handles
- Lots of overhead
- Every process has at least one kernel thread

■ Kernel Threads

- Shared virtual address space
- Contains running state data
- Less overhead
- From OS's point of view, this is what is scheduled to run on a CPU





Processes, User Threads, Kernel Threads

■ User Threads

- Shared virtual address space, contains running state data
- Kernel unaware
- Even less overhead





Trade-Offs

■ Processes

- Secure and isolated 😊
- Kernel aware 😊
- Creating a new process (address space!) brings lots of overhead ☹️





Trade-Offs (cont.)

■ Kernel Threads

- No need to create a new address space 😊
- No need to change address space in context switch 😊
- Kernel aware 😊
- Still need to enter kernel to context switch ☹️

■ User Threads

- No new address space, no need to change address space 😊
- No need to enter kernel to switch 😊
- Kernel is unaware; No multiprocessing; I/O blocks all user threads; ☹️





Thread Libraries

- **Thread library** Provides Programmer with API for Creating and Managing Threads
- Two Primary Ways of Implementing
 - Library entirely in **user space**
 - **Kernel-level** library supported by OS





Thread Libraries (cont.)

■ Library Entirely in User Space

- With no kernel support
- All code/data structures in user space
- Invoking a function in API results in a local function call in user space not a system call

■ Kernel-Level Library Supported by OS

- Code and data structures for library exist in kernel space
- Invoking a function in API → system call





Thread Libraries (cont.)

■ Three Main Thread Libraries in Use

■ POSIX Pthreads Library

- Thread extension of POSIX standard
- May be provided as either a user- or kernel-level

■ Win32 Thread Library

- Kernel-level library

■ Java Thread API


- Allows threads to be created & managed directly in Java programs
- Implemented using a thread library of host system since JVM runs on top of host system





Pthreads

- May be Provided either as **User-Level** or **Kernel-Level**
- A POSIX Standard (IEEE 1003.1c) API for Thread **Creation** and **Synchronization**
- ***Specification***, not ***Implementation***
- API Specifies Behavior of Thread Library, Implementation is up to Development of Library
- Common in UNIX OS (Solaris, Linux, Mac OS X)





Pthreads Example

```
- #include <pthread.h>
#include <stdio.h>

int sum; /* this data is shared by the thread(s) */
void *runner(void *param); /* threads call this function */

int main(int argc, char *argv[])
{
    pthread_t tid; /* the thread identifier */
    pthread_attr_t attr; /* set of thread attributes */

    if (argc != 2) {
        fprintf(stderr, "usage: a.out <integer value>\n");
        return -1;
    }
    if (atoi(argv[1]) < 0) {
        fprintf(stderr, "%d must be >= 0\n", atoi(argv[1]));
        return -1;
    }
}
```



Pthreads Example (Cont.)

```
/* get the default attributes */
pthread_attr_init(&attr);
/* create the thread */
pthread_create(&tid,&attr,runner,argv[1]);
/* wait for the thread to exit */
pthread_join(tid,NULL);

printf("sum = %d\n",sum);
}

/* The thread will begin control in this function */
void *runner(void *param)
{
    int i, upper = atoi(param);
    sum = 0;

    for (i = 1; i <= upper; i++)
        sum += i;

    pthread_exit(0);
}
```





Pthreads Code for Joining 10 Threads

```
#define NUM_THREADS 10
```

```
/* an array of threads to be joined upon */  
pthread_t workers[NUM_THREADS];
```

```
for (int i = 0; i < NUM_THREADS; i++)  
    pthread_join(workers[i], NULL);
```





Windows Multithreaded C Program

```
#include <windows.h>
#include <stdio.h>
DWORD Sum; /* data is shared by the thread(s) */

/* the thread runs in this separate function */
DWORD WINAPI Summation(LPVOID Param)
{
    DWORD Upper = *(DWORD*)Param;
    for (DWORD i = 0; i <= Upper; i++)
        Sum += i;
    return 0;
}

int main(int argc, char *argv[])
{
    DWORD ThreadId;
    HANDLE ThreadHandle;
    int Param;

    if (argc != 2) {
        fprintf(stderr, "An integer parameter is required\n");
        return -1;
    }
    Param = atoi(argv[1]);
    if (Param < 0) {
        fprintf(stderr, "An integer >= 0 is required\n");
        return -1;
    }
}
```





Windows Multithreaded C Program (Cont.)

```
/* create the thread */
ThreadHandle = CreateThread(
    NULL, /* default security attributes */
    0, /* default stack size */
    Summation, /* thread function */
    &Param, /* parameter to thread function */
    0, /* default creation flags */
    &ThreadId); /* returns the thread identifier */

if (ThreadHandle != NULL) {
    /* now wait for the thread to finish */
    WaitForSingleObject(ThreadHandle, INFINITE);

    /* close the thread handle */
    CloseHandle(ThreadHandle);

    printf("sum = %d\n", Sum);
}
```





Java Threads

- Java Threads Managed by JVM
- Typically Implemented using Threads Model provided by Underlying OS
- Each Java Program consists of at Least One Single Thread





Java Multithreaded Program

```
class Sum
{
    private int sum;

    public int getSum() {
        return sum;
    }

    public void setSum(int sum) {
        this.sum = sum;
    }
}

class Summation implements Runnable
{
    private int upper;
    private Sum sumValue;

    public Summation(int upper, Sum sumValue) {
        this.upper = upper;
        this.sumValue = sumValue;
    }

    public void run() {
        int sum = 0;
        for (int i = 0; i <= upper; i++)
            sum += i;
        sumValue.setSum(sum);
    }
}
```





Java Multithreaded Program (Cont.)

```
public class Driver
{
    public static void main(String[] args) {
        if (args.length > 0) {
            if (Integer.parseInt(args[0]) < 0)
                System.err.println(args[0] + " must be >= 0.");
            else {
                Sum sumObject = new Sum();
                int upper = Integer.parseInt(args[0]);
                Thread thrd = new Thread(new Summation(upper, sumObject));
                thrd.start();
                try {
                    thrd.join();
                    System.out.println
                        ("The sum of "+upper+" is "+sumObject.getSum());
                } catch (InterruptedException ie) { }
            }
        }
        else
            System.err.println("Usage: Summation <integer value>");
    }
}
```





Implicit Threading

- Growing in Popularity as Numbers of Threads Increase, Program Correctness more Difficult with Explicit Threads
- Creation and Management of Threads done by Compilers and Run-Time Libraries rather than Programmers
- Two Methods explored (Reading Assignment)
 - OpenMP
 - Grand Central Dispatch





Threading Issues

- Semantics of **fork()** and **exec()** system calls
- Signal Handling
 - Synchronous and asynchronous
- Thread Cancellation of Target Thread
 - Asynchronous or deferred
- Scheduler Activations





Semantics of `fork()` and `exec()`

■ Does `fork()` Duplicate only Calling Thread or all Threads?

- Some UNIXes have two versions of `fork`

■ What about `exec`?

- Usually works as normal
- Replaces running process including all threads





Signal Handling

- **Signals** Used in UNIX to Notify a Process that a Particular Event has Occurred
- **Signal Handler** Used to Deal with Signals
 1. Signal is generated by particular event
 2. Signal is delivered to a process
 3. Signal is handled by one of two signal handlers:
 - ❑ default
 - ❑ user-defined





Signal Handling (cont.)

- Every Signal has **Default Handler** that Kernel Runs when Handling Signal
 - **User-defined signal handler** can override default
 - For single-threaded, signal delivered to process





Signal Handling (cont.)

- Where should a Signal be Delivered for Multi-Threaded?
 - Deliver signal to thread to which signal applies
 - Deliver signal to **every thread in process**
 - Deliver signal to **certain threads in process**
 - **Assign a specific thread to receive all signals for process**





Thread Pools

■ Create a Number of Threads in a Pool where they Await Work

```
DWORD WINAPI PoolFunction(AVOID Param) {  
    /*  
     * this function runs as a separate thread.  
     */  
}
```

■ Advantages:

- Usually slightly faster to service a request with an existing thread than create a new thread
- Allows number of threads in application(s) to be bound to size of pool

■ Windows API supports thread pools





Thread Cancellation

- Terminating a Thread before it has finished
- Thread to be Canceled is **Target Thread**
- Two General Approaches
 - **Asynchronous cancellation** terminates target thread immediately
 - **Deferred cancellation** periodically checks for cancellation
- Pthread Code

```
pthread_t tid;  
  
/* create the thread */  
pthread_create(&tid, 0, worker, NULL);  
  
...  
  
/* cancel the thread */  
pthread_cancel(tid);
```





Thread Cancellation (Cont.)

- Invoking Thread Cancellation Requests
Cancellation, but Actual Cancellation Depends on Thread State

Mode	State	Type
Off	Disabled	–
Deferred	Enabled	Deferred
Asynchronous	Enabled	Asynchronous


- If Thread has Cancellation Disabled, Cancellation Remains Pending until Thread Enables it
- Default Type is Deferred
 - Cancellation only occurs when thread reaches **cancellation point**
 - ▶ I.e. `pthread_testcancel()`
 - ▶ Then **cleanup handler** is invoked
- On Linux, thread cancellation handled through signals





Scheduler Activations

- M:M/Two-Level Models Require communication to maintain appropriate number of kernel threads allocated to application



- **LightWeight Process (LWP)**

- Typically use an intermediate data structure between user and kernel threads
- Appears to be a virtual processor on which process can schedule user thread to run
- Each LWP attached to kernel thread
- How many LWPs to create?





Operating System Examples

■ Reading Assignments

- Windows threads
- Linux threads
- Upcalls





Windows Threads

- Windows implements Windows API
 - Primary API for Win98, WinNT, Win2000, WinXP, Win7
- Implements one-to-one Mapping, Kernel-level
- Each Thread Contains
 - A thread id
 - Register set representing state of processor
 - Separate user and kernel stacks for when thread runs in user mode or kernel mode
 - Private data storage area used by run-time libraries and dynamic link libraries (DLLs)
- **Context of Thread:** register set, stacks, and private storage area





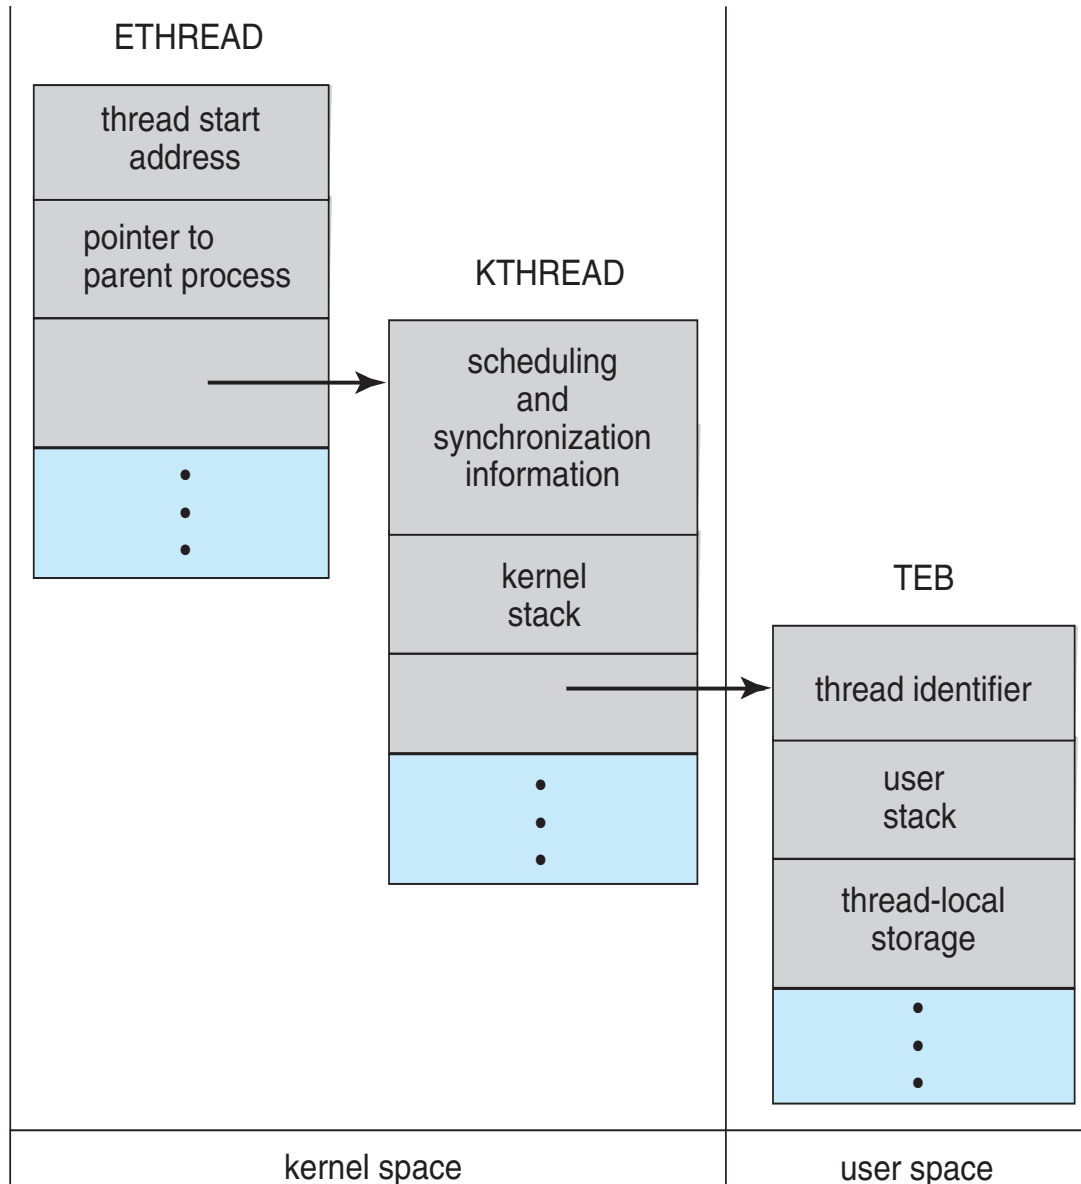
Windows Threads (Cont.)

- Primary Data Structures of a Thread include:
 - ETHREAD (executive thread block)
 - ▶ Includes pointer to process to which thread belongs and to KTHREAD, in kernel space
 - KTHREAD (kernel thread block)
 - ▶ Scheduling and synchronization info, kernel-mode stack, pointer to TEB, in kernel space
 - TEB (thread environment block)
 - ▶ Thread id, user-mode stack, thread-local storage, in user space





Windows Threads Data Structures





Linux Threads


- Linux refers to them as ***tasks*** rather than ***threads***
- Thread creation using `clone()` system call
- `clone()` allows a child task to share address space of parent task (process)
 - Flags control behavior
- `struct task_struct` points to process data structures (shared or unique)

flag	meaning
CLONE_FS	File-system information is shared.
CLONE_VM	The same memory space is shared.
CLONE_SIGHAND	Signal handlers are shared.
CLONE_FILES	The set of open files is shared.





Scheduler Activations (cont.)



■ Scheduler Activations Provides **upcalls**

- A communication mechanism from kernel to **upcall handler** in thread library
- This communication allows an application to maintain correct number kernel threads



End of Lecture 4

