


# Lecture 7: Deadlocks





# Lecture 7: Deadlocks

---

- System Model
- Deadlock Characterization
- Methods for Handling Deadlocks
- Deadlock Prevention
- Deadlock Avoidance
- Deadlock Detection
- Recovery from Deadlock





# Lecture Objectives

---

- To Develop a **Description of Deadlocks**, which Prevent sets of Concurrent Processes from Completing their Tasks
- To Present a number of Different Methods for **Preventing** or **Avoiding** Deadlocks in a Computer System





# System Model

## ■ System Consists of Resources

- To be distributed among a number of competing processes

## ■ Resource Types $R_1, R_2, \dots, R_m$

*E.g.,: CPU cycles, memory space, I/O devices*

## ■ Each Resource Type $R_i$ has $W_i$ instances

## ■ Each Process Utilizes a Resource as follows

### ● Request, Use, Release

- ▶ Accomplished using system calls such as request() and release() device or open() and close() file, OR allocate() and free() memory





# System Model (cont.)

---

## ■ Physical Resources

- Printers, tape drives, memory space, or CPU cores/cycles

## ■ Logical Resources

- Files, semaphores, and monitors

## ■ Example: Consider a system with 3 CD RW drives and 3 running processes

- Each process holds one CD RW drive and now requests another drive → Deadlock

## ■ Multithreaded Programs:

- Good candidate for deadlock
- Multiple threads compete for shared resources





# Deadlock Characterization

Deadlock can arise if four conditions hold simultaneously

- **Mutual Exclusion:** only one process at a time can use a resource
- **Hold and Wait:** a process holding at least one resource is waiting to acquire additional resources held by other processes
- **No Preemption:** a resource can be released only voluntarily by process holding it, after that process has completed its task





# Deadlock Characterization (cont.)

- **Circular Wait:** there exists a set  $\{P_0, P_1, \dots, P_n\}$  of waiting processes such that  $P_0$  is waiting for a resource that is held by  $P_1$ ,  $P_1$  is waiting for a resource that is held by  $P_2$ , ...,  $P_{n-1}$  is waiting for a resource that is held by  $P_n$ , and  $P_n$  is waiting for a resource that is held by  $P_0$ .





# Resource-Allocation Graph

A set of vertices  $V$  and a set of edges  $E$ .

■  $V$  is partitioned into two types:

- $P = \{P_1, P_2, \dots, P_n\}$ , set consisting of all processes in the system
- $R = \{R_1, R_2, \dots, R_m\}$ , set consisting of all resource types in the system

■ **Request Edge** – directed edge  $P_i \rightarrow R_j$

■ **Assignment Edge** – directed edge  $R_j \rightarrow P_i$








# Resource-Allocation Graph (cont.)


## ■ Process




## ■ Resource Type with 4 instances




## ■ $P_i$ requests instance of $R_j$



## ■ $P_i$ is holding an instance of $R_j$



# Example of a Resource Allocation Graph



■  $P = \{P_1, P_2, P_3\}$

■  $R = \{R_1, R_2, R_3, R_4\}$


●  $R_1$ : one instance,  $R_2$  two instances,  $R_3$  one instance, and  $R_4$  three instances

■  $E = \{P_1 \rightarrow R_1, P_2 \rightarrow R_3, \dots, R_3 \rightarrow P_3\}$






# Resource Allocation Graph With a Deadlock





# Graph With A Cycle But No Deadlock





# Basic Facts

- If Graph Contains **no Cycles**  $\Rightarrow$  **No Deadlock**
- If Graph Contains a Cycle  $\Rightarrow$ 
  - Only one instance per resource type  $\rightarrow$  **deadlock**
  - Cycle involves only a set of resource types, each of which has only a single instance  $\rightarrow$  **deadlock**
    - ▶ In the above two cases, cycle is necessary & sufficient condition for existence of deadlock
  - Several instances per resource type  $\rightarrow$  **possibility of deadlock**
    - ▶ In this case, cycle is necessary but not sufficient





# Methods for Handling Deadlocks

- Ensure that System will **Never** Enter a Deadlock State
  - Deadlock **prevention**: try to **violate** one of necessary **conditions** for deadlock
  - Deadlock **avoidance**: try to **regulate** how/when requests can be made to acquire resources
    - ▶ More **conservative** approach than deadlock prevention
- Allow System to **enter a Deadlock State** and then **recover**





# Methods for Handling Deadlocks (cont.)

- Ignore Problem and Pretend that Deadlocks Never occur in system
  - Used by most OSes, including UNIX
  - Up to application developer to detect and handle deadlocks
- What if Deadlocks are not Resolved?
  - Deterioration of system performance
    - ▶ Eventually need a manual restart
  - Deadlock occur very infrequent → cheaper approach in mainstream applications
    - ▶ Instead of employing prevention, avoidance, or detection and recovery methods





# Deadlock Prevention

---

## ■ Mutual Exclusion

- Not required for sharable resources
- A process never needs to wait for a sharable resources
- Must hold for non-sharable resources
- Example
  - ▶ Read-only files







# Deadlock Prevention (cont.)

- **Hold and Wait** – must guarantee that whenever a process requests a resource, it does not hold any other resources
  - **Solution 1**: Require process to request and be allocated **all its resources** before it **begins** execution
  - **Solution 2**: Or allow process to request resources only when process **has none** allocated to it
  - **Cons**
    - ▶ Low resource utilization ☹️
    - ▶ Starvation possible ☹️





# Deadlock Prevention (cont.)

## ■ No Preemption

- If a process that is holding some resources requests another resource that cannot be immediately allocated to it, then **all resources currently** being held are **released**
- **Preempted resources** are added to list of resources for which the process is **waiting**
- Process will be **restarted only** when it can **regain** its **old** resources, **as well as new ones**
- This protocol **applicable only** to resources whose state can be easily saved and restored later
  - ▶ CPU registers and memory space: applicable 😊
  - ▶ Printers and tape drives: not (easily) applicable ☹️






# Deadlock Prevention (cont.)

■ **Circular Wait** – impose a **total ordering** of all resource types, and require that each process requests resources in an **increasing order** of enumeration

- A process which holds  $R(i)$ , can request instance of  $R(j)$  if  $F(Rj) > F(Ri)$
- Ensuring order is possibility of application developer
  - ▶ Can use **lock-order verifier** (e.g., **witness** in FreeBSD)

## ■ Example

- $F(\text{tape})=1$ ,  $F(\text{disk drive})=5$ , and  $F(\text{printer})=12$





# Deadlock Example

```
/* thread one runs in this function */
void *do_work_one(void *param)
{
    pthread_mutex_lock(&first_mutex);
    pthread_mutex_lock(&second_mutex);
    /** * Do some work */
    pthread_mutex_unlock(&second_mutex);
    pthread_mutex_unlock(&first_mutex);
    pthread_exit(0);
}

/*****/
/* thread two runs in this function */
void *do_work_two(void *param)
{
    pthread_mutex_lock(&second_mutex);
    pthread_mutex_lock(&first_mutex);
    /** * Do some work */
    pthread_mutex_unlock(&first_mutex);
    pthread_mutex_unlock(&second_mutex);
    pthread_exit(0);
}
```





# Deadlock Example with Lock Ordering

## ■ Lock Ordering does not Guarantee Deadlock Prevention if Locks can be acquired Dynamically

- Ordering is broken with **unordered arguments**

```
void transaction(Account from, Account to, double amount)
{
    mutex lock1, lock2;
    lock1 = get_lock(from);
    lock2 = get_lock(to);
    acquire(lock1);
        acquire(lock2);
            withdraw(from, amount);
            deposit(to, amount);
        release(lock2);
    release(lock1);
}
```

- Transactions 1 and 2 execute concurrently.
- Transaction 1 transfers \$25 from account A to account B, &  
Transaction 2 transfers \$50 from account B to account A





# Deadlock Avoidance

Requires that system has some additional *a priori* information available

- Simplest and most useful model requires that each process declare *maximum number* of resources of each type that it may need
- Deadlock-avoidance algorithm dynamically examines *resource-allocation state* to ensure that there can never be a *circular-wait* condition
- “Resource-allocation *state*” is defined by number of *available* and *allocated* resources, and *maximum demands* of processes





# Safe State

- When a process requests an available resource, system must decide if immediate allocation **leaves** system in a **safe state**
- System is in **safe state** if there exists a sequence  $\langle P_1, P_2, \dots, P_n \rangle$  (aka, **safe sequence**) of ALL processes in systems such that for each  $P_i$ , resources that  $P_i$  can still request can be satisfied by **currently available resources + resources held by all  $P_j$ , with  $j < i$**





# Safe State (cont.)

---

## ■ That is:

- If  $P_i$  resource needs are not immediately available, then  $P_i$  can wait until all  $P_j$  have finished
- When  $P_j$  is finished,  $P_i$  can obtain needed resources, execute, return allocated resources, and terminate
- When  $P_i$  terminates,  $P_{i+1}$  can obtain its needed resources, and so on







# Basic Facts


---

- If a System is in **Safe State**  $\Rightarrow$  **No Deadlocks**
- If a System is in **Unsafe State**  $\Rightarrow$  **Possibility of Deadlock**
- **Avoidance**  $\Rightarrow$  **Ensure that a System will Never enter an Unsafe State**





# Safe, Unsafe, Deadlock State





# Avoidance Algorithms

---

- **Single Instance** of a Resource Type
  - Use a resource-allocation graph
  
- **Multiple Instances** of a Resource Type
  - Use banker's algorithm
    - ▶ Reading assignment





# Resource-Allocation Graph Scheme

## ■ Claim edge

- $P_i \rightarrow R_j$  indicates that process  $P_i$  may request resource  $R_j$  (represented by a dashed line)

## ■ Claim edge converts to request edge when a process requests a resource

## ■ Request edge converted to an assignment edge when resource is allocated to process


## ■ Resource is Released by a process → Assignment Edge reconverts to a claim edge

## ■ Resources must be claimed *a priori* in system






# Resource-Allocation Graph



# Unsafe State In Resource-Allocation Graph



Why not assigning  $R_2$  to  $P_2$  after making sure that  $P_1$  doesn't need  $R_2$  in the near future?





# Resource-Allocation Graph Algorithm

---

- Suppose that  $P_i$  requests a  $R_j$
- Request can be Granted only if Converting Request Edge to an Assignment Edge does not Result in Formation of a Cycle in Resource Allocation Graph
- Algorithm Complexity  $O(n^2)$
- Low Resource Utilization
  - Resources might be available and not be allocated to processes





# Deadlock Detection

---

■ Allow System to Enter Deadlock State

■ Detection Algorithm

- Single instance of each resource type
- Multiple instances of a resource type
  - ▶ Reading assignment

■ Recovery Scheme







# Single Instance of Each Resource Type

---

## ■ Maintain **wait-for** graph

- Nodes are processes
- $P_i \rightarrow P_j$  if  $P_i$  is waiting for  $P_j$

## ■ **Periodically** invoke an Algorithm that Searches for a **Cycle** in graph

- If there is a cycle  $\rightarrow$  there exists a deadlock


## ■ An Algorithm to **Detect** a Cycle in a Graph Requires an **Order** of $n^2$ operations

- Where  $n$  is number of vertices in graph






# Resource-Allocation Graph and Wait-for Graph



(a)

**Resource-Allocation  
Graph**



(b)

**Corresponding  
wait-for Graph**





# Detection-Algorithm Usage

- **When**, and **How often**, to Invoke depends on:
  - How often a deadlock is likely to occur?
  - How many processes will need to be rolled back?
    - ▶ One for each disjoint cycle
- If **Detection Algorithm** is Invoked Arbitrarily, there may be many cycles in resource graph and so we would **not be able to tell** which of many deadlocked processes “**caused**” **deadlock**





# Recovery from Deadlock: **Process Termination**

---

## ■ **Abort** all Deadlocked Processes

- Significant expense
- Most processes need to be restarted for re-computation

## ■ Abort one process at a time until Deadlock Cycle is Eliminated

- After each process is aborted, a deadlock-detection algorithm must be invoked
  - ▶ To see if any process is still deadlocked





# Recovery from Deadlock: **Process Termination**

---

## ■ Issues with Aborting a Process

- Updating a file
- Printing data

## ■ In which Order should we Choose to Abort?

1. Priority of process
2. How long process has computed, and how much longer to completion
3. Resources process has used
4. Resources process needs to complete
5. How many processes will need to be terminated
6. Is process interactive or batch?





# Recovery from Deadlock: Resource Preemption

---

- **Selecting a Victim** – minimize cost
- **Rollback** – return to some safe state, restart process for that state
- **Starvation** – same process may always be picked as victim, include number of rollback in cost factor





# Reading Assignment

---

- Banker's Algorithm
- Deadlock Detection
  - Multiple instances of a resource type





# Banker's Algorithm

---

- Multiple Instances
- Each Process must a Priori Claim Max Use
- When a Process Requests a Resource it may have to Wait
- When a Process Gets all its Resources it must Return them in a Finite Amount of Time







# Data Structures for Banker's Algorithm

Let  $n$  = number of processes, and  $m$  = number of resources types.

- **Available:** Vector of length  $m$ . If available  $[j] = k$ , there are  $k$  instances of resource type  $R_j$  available
- **Max:**  $n \times m$  matrix. If  $Max[i, j] = k$ , then process  $P_i$  may request at most  $k$  instances of resource type  $R_j$
- **Allocation:**  $n \times m$  matrix. If  $Allocation[i, j] = k$  then  $P_i$  is currently allocated  $k$  instances of  $R_j$
- **Need:**  $n \times m$  matrix. If  $Need[i, j] = k$ , then  $P_i$  may need  $k$  more instances of  $R_j$  to complete its task  
 $Need[i, j] = Max[i, j] - Allocation[i, j]$





# Safety Algorithm

1. Let **Work** and **Finish** be vectors of length  $m$  and  $n$ , respectively. Initialize:

**Work** = **Available**

**Finish** [ $i$ ] = **false** for  $i = 0, 1, \dots, n-1$

2. Find an  $i$  such that both:

(a) **Finish** [ $i$ ] = **false**

(b) **Need** <sub>$i$</sub>  ≤ **Work**

If no such  $i$  exists, go to step 4

3. **Work** = **Work** + **Allocation** <sub>$i$</sub>

**Finish** [ $i$ ] = **true**

go to step 2

4. If **Finish** [ $i$ ] == **true** for all  $i$ , then the system is in a safe state





# Resource-Request Algorithm for Process $P_i$

**$Request_i$**  = request vector for process  $P_i$ .

If  **$Request_i[j] = k \rightarrow P_i$**  wants  $k$  instances of  $R_j$

1. If  **$Request_i \leq Need_i$**  go to step 2. Otherwise, raise error condition (as process has exceeded its maximum claim)
2. If  **$Request_i \leq Available$** , go to step 3. Otherwise  $P_i$  must wait, since resources are not available
3. Pretend to allocate requested resources to  $P_i$  by modifying the state as follows:

**$Available = Available - Request_i;$**

**$Allocation_i = Allocation_i + Request_i;$**

**$Need_i = Need_i - Request_i;$**

- If safe  $\Rightarrow$  the resources are allocated to  $P_i$
- If unsafe  $\Rightarrow P_i$  must wait, and old resource-allocation state is restored





# Example of Banker's Algorithm

■ 5 processes  $P_0$  through  $P_4$ ;

3 resource types:

$A$  (10 instances),  $B$  (5 instances), and  $C$  (7 instances)

■ Snapshot at time  $T_0$ :

	<u>Allocation</u>	<u>Max</u>	<u>Available</u>
	$A \ B \ C$	$A \ B \ C$	$A \ B \ C$
$P_0$	0 1 0	7 5 3	3 3 2
$P_1$	2 0 0	3 2 2	
$P_2$	3 0 2	9 0 2	
$P_3$	2 1 1	2 2 2	
$P_4$	0 0 2	4 3 3	





## Example (cont.)

- Content of ***Need*** is defined to be ***Max – Allocation***

*Need*

*A B C*

*P*<sub>0</sub> 7 4 3

*P*<sub>1</sub> 1 2 2

*P*<sub>2</sub> 6 0 0

*P*<sub>3</sub> 0 1 1

*P*<sub>4</sub> 4 3 1

- System is in a **Safe** State

- Since sequence  $\langle P_1, P_3, P_4, P_2, P_0 \rangle$  satisfies safety criteria





## Example: $P_1$ Request (1,0,2)

- Check that Request  $\leq$  Available (i.e.,  $(1,0,2) \leq (3,3,2) \Rightarrow$  true

	<u>Allocation</u>			<u>Need</u>			<u>Available</u>		
	A	B	C	A	B	C	A	B	C
$P_0$	0	1	0	7	4	3	2	3	0
$P_1$	3	0	2	0	2	0			
$P_2$	3	0	2	6	0	0			
$P_3$	2	1	1	0	1	1			
$P_4$	0	0	2	4	3	1			

- Executing safety algorithm shows that sequence  $\langle P_1, P_3, P_4, P_0, P_2 \rangle$  satisfies safety requirement
- Can request for (3,3,0) by  $P_4$  be granted?
- Can request for (0,2,0) by  $P_0$  be granted?





# Several Instances of a Resource Type

- **Available:** A vector of length  $m$  indicates number of available resources of each type
- **Allocation:** An  $n \times m$  matrix defines number of resources of each type currently allocated to each process
- **Request:** An  $n \times m$  matrix indicates current request of each process
  - If  $Request[i][j] = k$ ,  $\rightarrow$  process  $P_i$  is requesting  $k$  more instances of resource type  $R_j$





# Detection Algorithm

1. Let ***Work*** and ***Finish*** be vectors of length ***m*** and ***n***, respectively Initialize:

(a) ***Work = Available***

(b) For  $i = 1, 2, \dots, n$ , if ***Allocation<sub>i</sub>  $\neq 0$*** , then  
***Finish[i] = false***; otherwise, ***Finish[i] = true***

2. Find an index ***i*** such that both:

(a) ***Finish[i] == false***

(b) ***Request<sub>i</sub>  $\leq$  Work***

If no such ***i*** exists, go to step 4







# Detection Algorithm (cont.)

3.  **$Work = Work + Allocation_i$**   
 **$Finish[i] = true$**   
go to step 2

4. If  **$Finish[i] == false$** , for some  $i$ ,  $1 \leq i \leq n$ ,  
then the system is in deadlock state.  
Moreover, if  **$Finish[i] == false$** , then  $P_i$  is  
deadlocked

**Algorithm requires an order of  $O(m \times n^2)$  operations  
to detect whether the system is in deadlocked state**





# Example of Detection Algorithm

- Five Processes  $P_0$  through  $P_4$ ; three Resource types  
A (7 instances), B (2 instances), and C (6 instances)

- Snapshot at time  $T_0$ :

	<u>Allocation</u>			<u>Request</u>			<u>Available</u>		
	A	B	C	A	B	C	A	B	C
$P_0$	0	1	0	0	0	0	0	0	0
$P_1$	2	0	0	2	0	2			
$P_2$	3	0	3	0	0	0			
$P_3$	2	1	1	1	0	0			
$P_4$	0	0	2	0	0	2			

- Sequence  $\langle P_0, P_2, P_3, P_1, P_4 \rangle$  will result in  **$Finish[i] = true$**  for all  $i$





# Example (cont.)


- $P_2$  requests an additional instance of type  $C$

Request

	$A$	$B$	$C$
$P_0$	0	0	0
$P_1$	2	0	2
$P_2$	0	0	1
$P_3$	1	0	0
$P_4$	0	0	2

- State of system?

- Can reclaim resources held by process  $P_0$ , but insufficient resources to fulfill other processes; requests
- Deadlock exists, consisting of processes  $P_1$ ,  $P_2$ ,  $P_3$ , and  $P_4$



# End of Lecture 7

---

