

Lecture 8: Main Memory





Lecture 8: Memory Management

- Background
- Swapping
- Contiguous Memory Allocation
- Segmentation
- Paging
- Structure of Page Table
- Example: Intel 32 and 64-bit Architectures
- Example: ARM Architecture





Objectives

- To Provide a Detailed Description of Various Ways of Organizing Memory Hardware
- To Discuss Various Memory-Management Techniques, including Paging and Segmentation
- To Provide a Description of Intel Pentium, which Supports both Pure Segmentation and Segmentation with Paging





Background

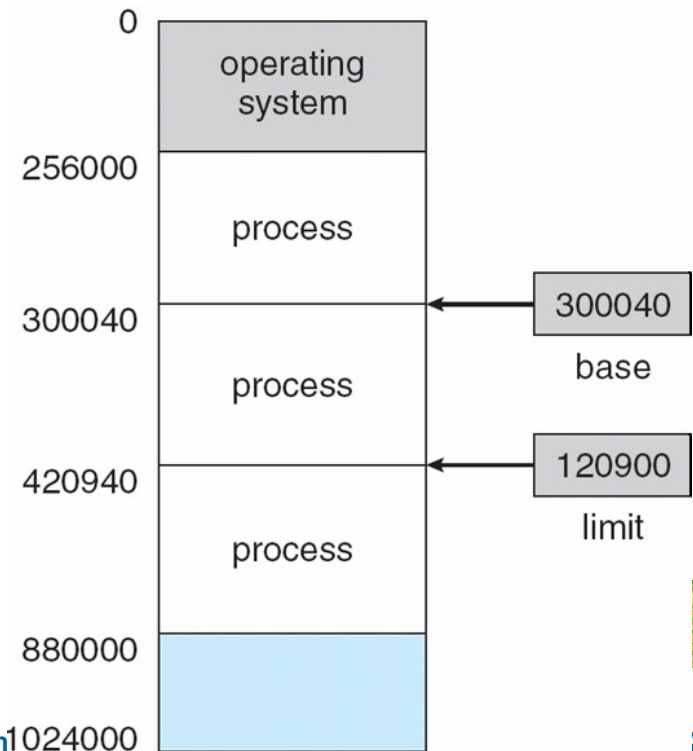
- Program must be **Brought** (from disk) into Memory and Placed within a Process for it to be Run
- **Main Memory (MM)** and **Registers** are **Only** Storage **CPU** can **Access Directly**
- Memory Unit **only Sees**
 - A Stream of **Addresses** + **Read** requests, or **Address** + **Data** and **Write** Requests
- **Register** Access in **one CPU Clock** (or Less)
- **MM** can Take **Many Cycles**, Causing a **Stall**
- **Cache** Sits between MM and CPU registers
- **Protection** of Memory required to **Ensure** Correct Operation





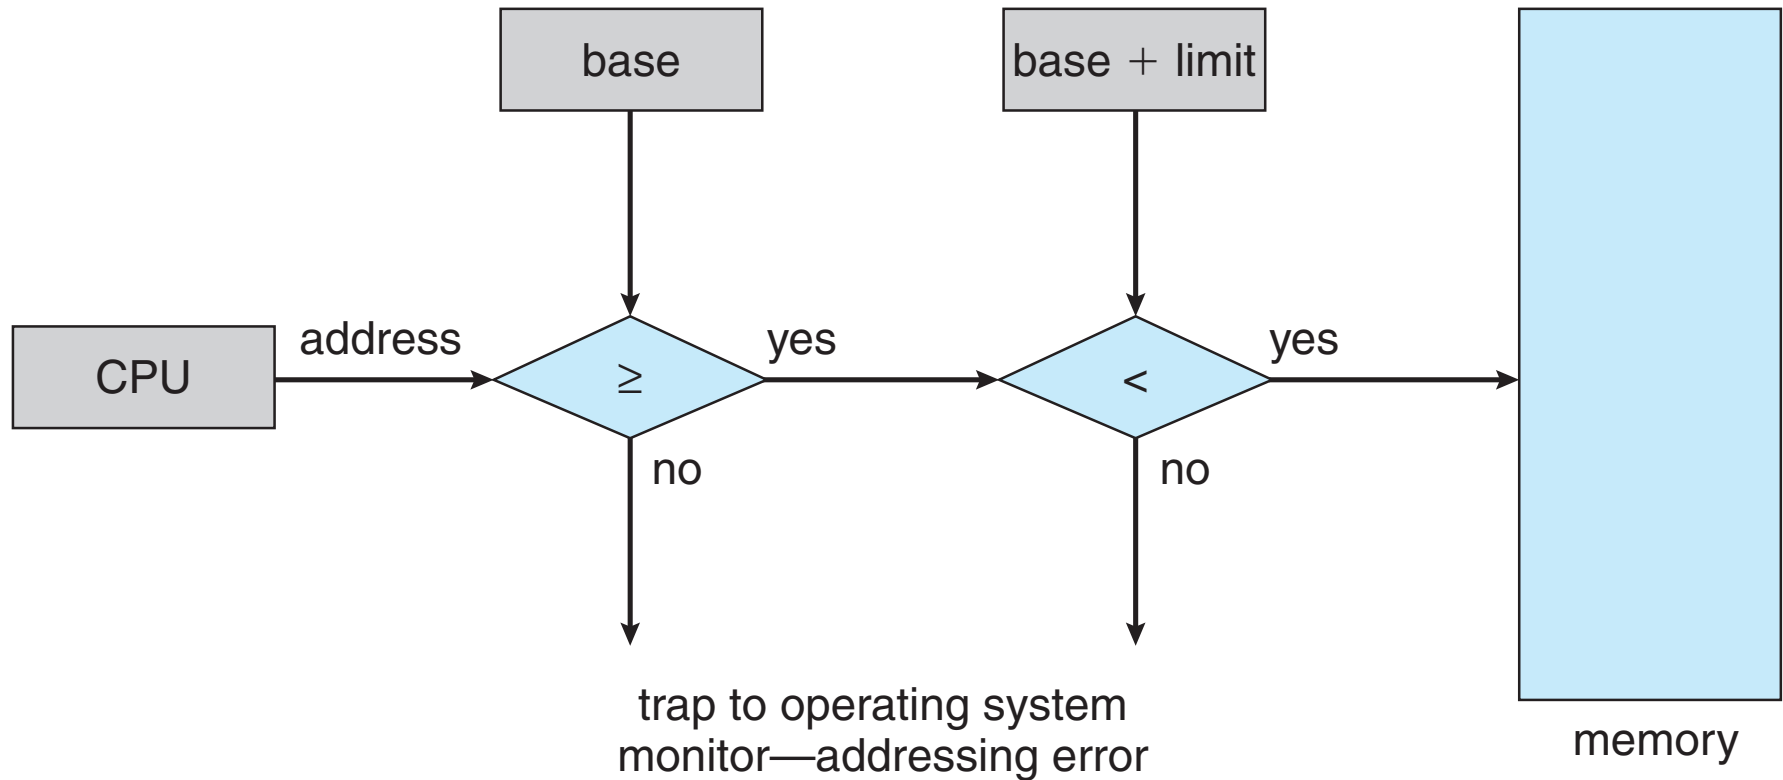
Base and Limit Registers

- A Pair of **Base** and **Limit Registers** Define Logical Address Space
- CPU must Check Every Memory Access
 - Generated in **user mode** to be sure it is between base and limit for that user





Hardware Address Protection





Address Binding

- Programs on Disk, Ready to be Brought into Memory to Execute Form an **Input Queue**
 - Without support, must be loaded into addr: **0000**
- Inconvenient to Have **First User Process** Physical Address Always at **0000**
 - How can it not be?
- Further, **Addresses** Represented in **Different Ways** at Different Stages of a Program's Life
 - Source code, compiled code, linker/loader addresses





Address Binding (cont.)

■ Different Ways of Addresses Representation

- Source code addresses usually **symbolic**
- Compiled code addresses **bind** to relocatable addresses
 - ▶ i.e. “14 bytes from beginning of module”
- Linker or loader will bind **relocatable** addresses to **absolute** addresses
 - ▶ i.e. 74014
- Each **binding** maps **one** address space to another





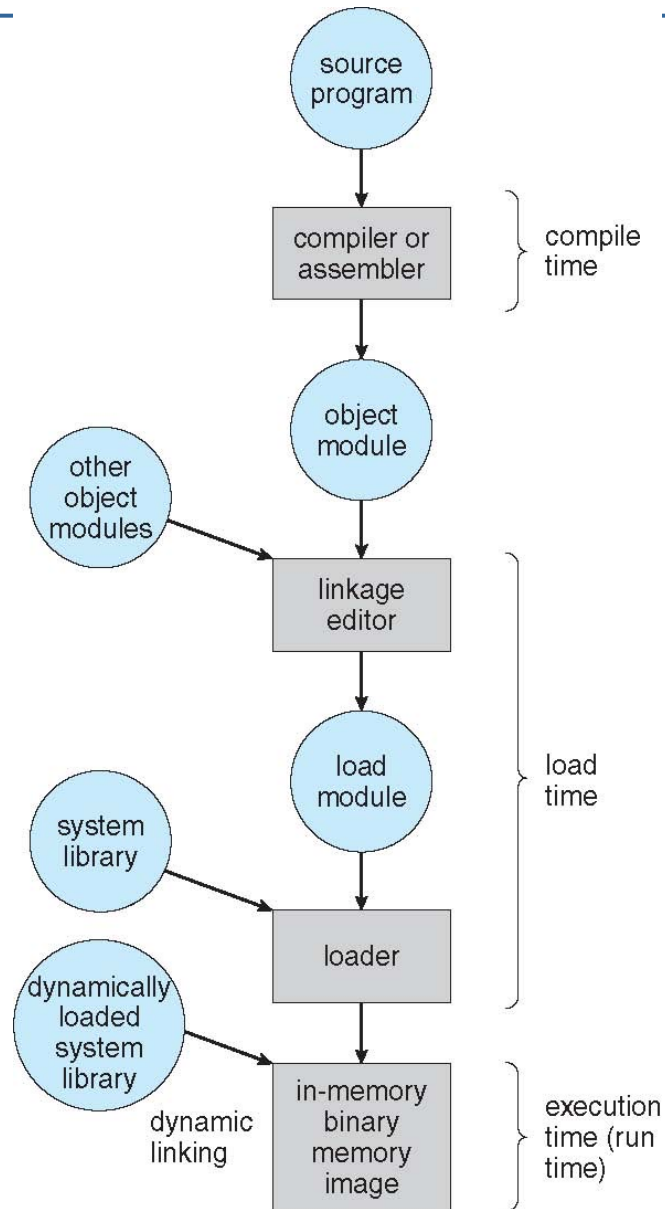
Binding of Instructions and Data to Memory


- Address Binding to Memory Addresses can Happen at Three Different Stages
 - **Compile time:** If memory location known a priori, **absolute code** can be generated
 - ▶ Must recompile code if starting location changes
 - **Load time:** Must generate **relocatable code** if memory location is not known at compile time
 - **Execution time:** Binding delayed until run time if process can be moved during its execution from one memory segment to another
 - ▶ Need HW support for address maps (e.g., base and limit registers)





Multistep Processing of a User Program





Logical vs. Physical Address Space

■ Concept of a Logical Address Space

- Central to Proper Memory Management
- Bound to a Separate **Physical Address Space**

■ **Logical address**

- Generated by CPU
- Aka, **virtual address**

■ **Physical Address**

- Address seen by memory unit

■ Logical and Physical Addresses

- Are **same** in **compile-time** and **load-time** address-binding schemes
- Differ in **execution-time** address-binding **scheme**



Logical vs. Physical Address Space (cont.)

■ Logical Address Space

- Set of all logical addresses generated by a program

■ Physical Address Space

- Set of all physical addresses generated by a program



Logical vs. Physical Address Space (cont.)

■ Memory Management Unit (MMU)

- HW device that at run-time maps **virtual** to **physical** address
- Many methods possible
 - ▶ Will be discussed in this lecture

■ To start, Consider Simple Scheme


- Value in **base register** is **added** to **every** address generated by a user process at time it is sent to memory
- Base register now called **relocation register**
- MS-DOS on Intel 80x86 used 4 relocation registers



Logical vs. Physical Address Space (cont.)

■ User Program Deals with *Logical* Addresses

- It never sees *real/physical* addresses
- Execution-time binding occurs when reference is made to location in memory
- Logical address bound to physical addresses





Dynamic Relocation using a Relocation Register

- Routine **not Loaded** until it is **Called**
- Better Memory-Space Utilization
 - Unused routine never loaded
- All Routines Kept on Disk in Relocatable Load Format
- **Useful** when **Large Amounts** of **Code** are Needed to Handle Infrequently Occurring Cases
- No Special Support from OS required
 - Implemented through program design
 - OS can help by providing libraries to implement dynamic loading





Dynamic Linking

- **Static Linking** – system libraries and program code combined by loader into binary program image
 - Wastes both disk space and main memory
- **Dynamic Linking** – linking postponed until execution time
 - Small piece of code, **Stub**, used to locate appropriate memory-resident library routine
 - Stub replaces itself with address of routine, and executes routine





Dynamic Linking (cont.)

- OS Checks if Routine is in Memory
 - If not in address space, add to address space
- Dynamic Linking is Particularly Useful for Libraries
 - Processes that use a language library execute **only one copy** of library code
 - Libraries updates will be automatically applied





Swapping

- A Process can be **Swapped Temporarily** out of Memory to a Backing Store, and Brought back into Memory for Continued Execution
- Why Swap out?
 - Quantum of round robin expired → need to bring a new process → not enough memory space
 - ▶ Total physical memory space of processes can exceed physical memory
- **Backing store**
 - Fast disk large enough to accommodate copies of all memory images for all users






Swapping (cont.)

- Major part of **Swap Time** is **Transfer Time**
 - Total transfer time is directly **proportional** to **amount of memory swapped**
- System Maintains a **Ready Queue** of ready-to-run processes
 - Which have memory images on disk (or in main memory)





Schematic View of Swapping





Context Switch Time and Swapping

- Next Processes to be Put on CPU Not in MM
 - Need to **swap out** a process & **swap in** target process
- Context **Switch Time** can then be **Very High**
- Example: 100MB Process
 - Swapping to HDD with transfer rate of 50MB/sec
 - Swap out time of 2000 ms
 - Plus swap in of same sized process
 - Total context switch swapping component time of 4000ms (4 seconds)





Swapping (cont.)

- **Q:** Does Swapped out Process Need to Swap Back into Same Physical Addresses?
- **Ans:** Depends on address binding method
- Standard Swapping is Too Time-Consuming
- ➔ Modified Versions of Swapping Found on Many Systems (i.e., UNIX, Linux, and Windows)
 - Swapping normally disabled
 - **Started** if more than threshold amount of memory allocated
 - **Disabled again** once memory demand reduced below threshold





Context Switch Time and Swapping (cont.)

■ Can reduce Context Switch Time

- If **reduce size of memory** swapped
 - ▶ by knowing how much memory really being used
- System calls to inform OS of memory use via `request_memory()` and `release_memory()`

■ Other **Constraints** as well on Swapping

- Pending I/O – can't swap out as I/O would occur to wrong process
- Or always transfer I/O to kernel space, then to I/O device
 - ▶ Known as **double buffering**, adds overhead





Memory Management Schemes

■ Main Memory must support OS + User Processes

- Limited Resource → must allocate efficiently

■ Memory Management Schemes

- Contiguous Allocation
- Segmentation
- Paging
- Segmentation + Paging





Contiguous Allocation

- Contiguous Allocation is one Early Method
- Main Memory Usually into two **partitions**:
 - Resident OS, usually held in low memory with interrupt vector
 - User processes then held in high memory
 - Each process contained in single contiguous section of memory





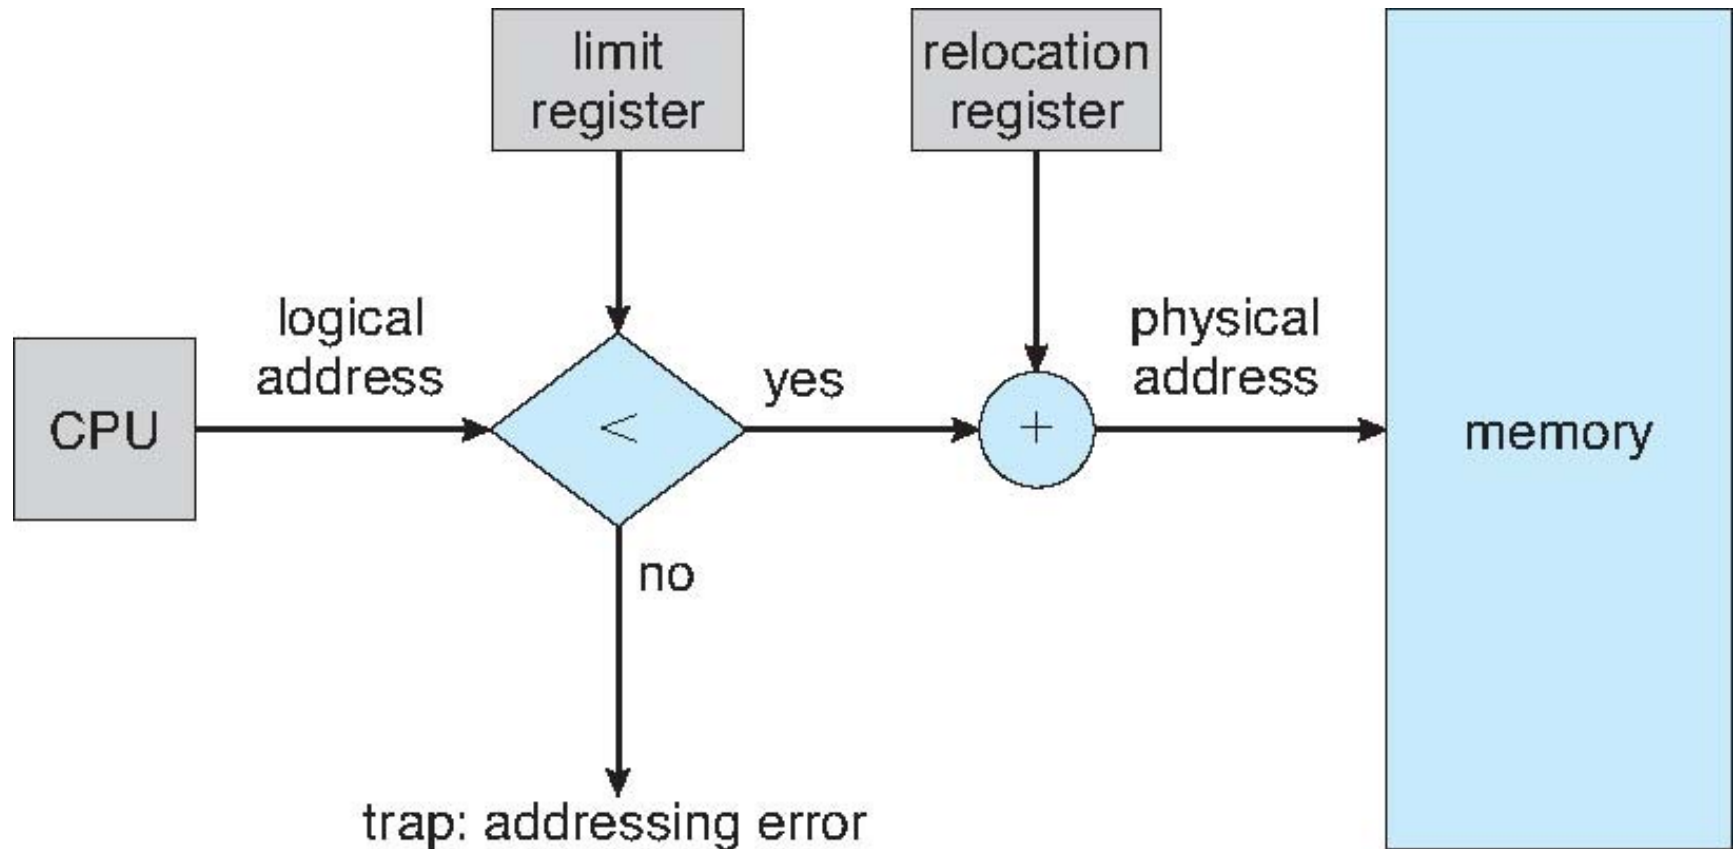
Contiguous Allocation (cont.)

- Relocation Registers used to Protect User Processes from each other, and from Changing OS code and data
 - Base register contains value of smallest physical address
 - Limit register contains range of logical addresses – each logical address must be less than the limit register
 - MMU maps logical address *dynamically*
 - Can then allow actions such as kernel code being **transient** and kernel changing size





HW Support for Relocation and Limit Registers





Multiple-Partition Allocation

■ Fixed-Sized Partitions

- Each partition only one process
- Degree of multiprogramming limited by number of partitions
- Originally used in IBM OS/360
- No longer in use





Multiple-Partition Allocation (cont.)

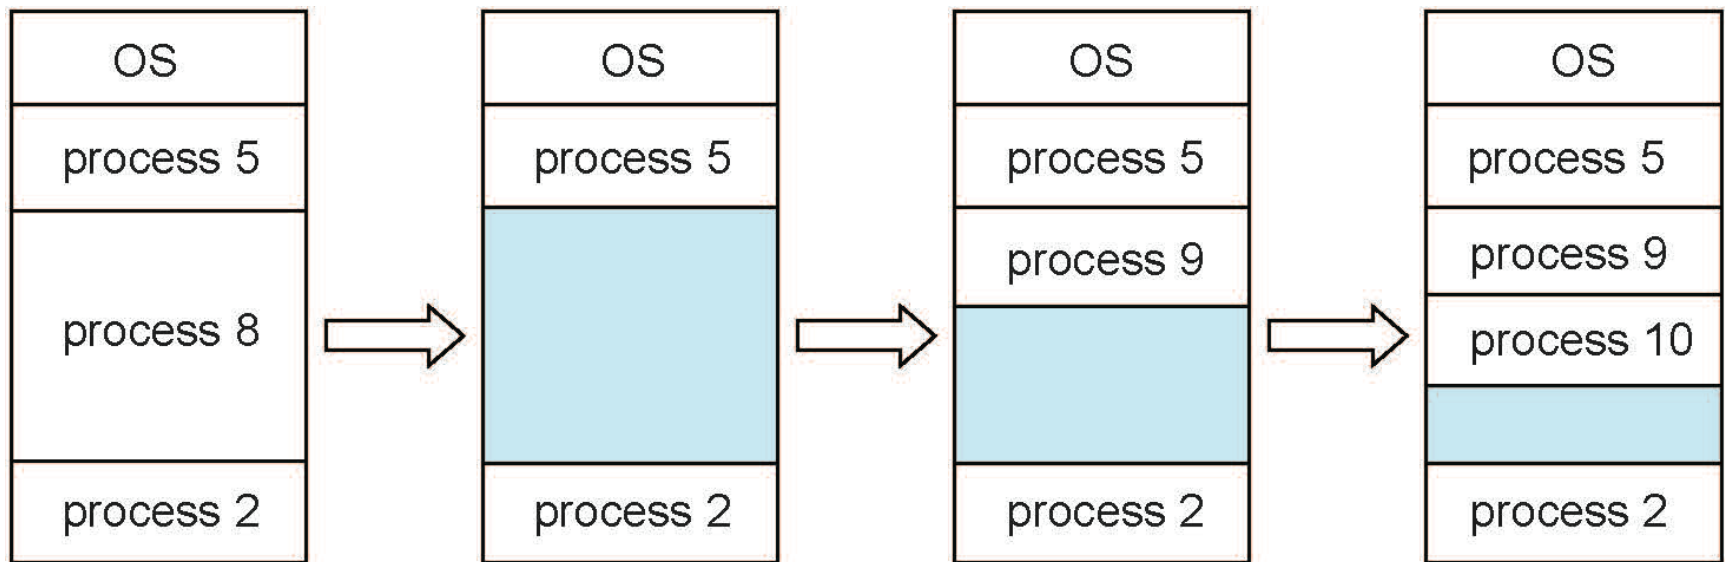
■ Variable Partition Scheme

- Sized to a given process' needs
- **Hole** – block of available memory
 - ▶ Holes of various size are scattered throughout memory
- When a process arrives, it is allocated memory from a hole large enough to accommodate it
- Process exiting frees its partition
 - ▶ **Adjacent free** partitions **combined**
- OS maintains information about:
 - a) **allocated** partitions b) **free** partitions (**hole**)





Multiple-Partition Allocation (cont.)





Dynamic Storage-Allocation Problem

How to satisfy a request of size n from a **list of free holes**?

- **First-Fit:** Allocate ***first*** hole that is big enough
 - **Best-Fit:** Allocate ***smallest*** hole that is big enough; must search **entire list**, unless ordered by size
 - Produces **smallest leftover hole**
 - **Worst-Fit:** Allocate ***largest*** hole; must also search **entire list**
 - Produces **largest leftover hole**
- First-fit and best-fit better than worst-fit in terms of speed and storage utilization**





Fragmentation

- **External Fragmentation** – Total memory space exists to satisfy a request, but it is not contiguous
- **Internal Fragmentation** – Allocated memory may be slightly larger than requested memory; this size difference is memory internal to a partition, but not being used
- First fit analysis reveals that given N blocks allocated, $0.5 N$ blocks lost to fragmentation
 - $1/3$ may be unusable -> **50-percent rule**





Fragmentation (cont.)

■ Solutions to External Fragmentation

● **Compaction**

- ▶ Shuffle memory contents to place all free memory together in one large block
 - E.g., move all used blocks to one end of memory
- ▶ Compaction is **possible only** if **relocation** is **dynamic**, and is done at execution time

● **Non-contiguous** memory allocation scheme

- ▶ Segmentation
- ▶ Paging

■ Backing store same Fragmentation Problem





Segmentation

- Memory-Management Scheme that Supports User View of Memory
 - Maps **programmer's** view to **physical** memory
- A Program is a Collection of Segments
 - A Segment is a Logical Unit such as:

main program
procedure
function
method
object
local variables
global variables
common block
stack
symbol table
arrays




Code
Global variables
Heap
Stack
Standard C library






User's View of a Program






Logical View of Segmentation



User Space



Physical Memory Space





Segmentation Architecture

- Logical Address consists of a Two Tuple:
 $\langle \text{segment-number}, \text{offset} \rangle$,
- **Segment table** – maps two-dimensional physical addresses; each table entry has:
 - **base** – contains **starting** physical address where segments reside in memory
 - **limit** – specifies **length** of segment
- **Segment-table base register (STBR)** points to segment table's location in memory
- **Segment-table length register (STLR)** indicates number of segments used by a program;
segment number **s** is legal if **s < STLR**





Segmentation Architecture (cont.)

■ Protection

- With each entry in segment table associate:
 - ▶ validation bit = 0 \Rightarrow illegal segment
 - ▶ read/write/execute privileges

■ Protection Bits Associated with Segments

- Code sharing occurs at segment level

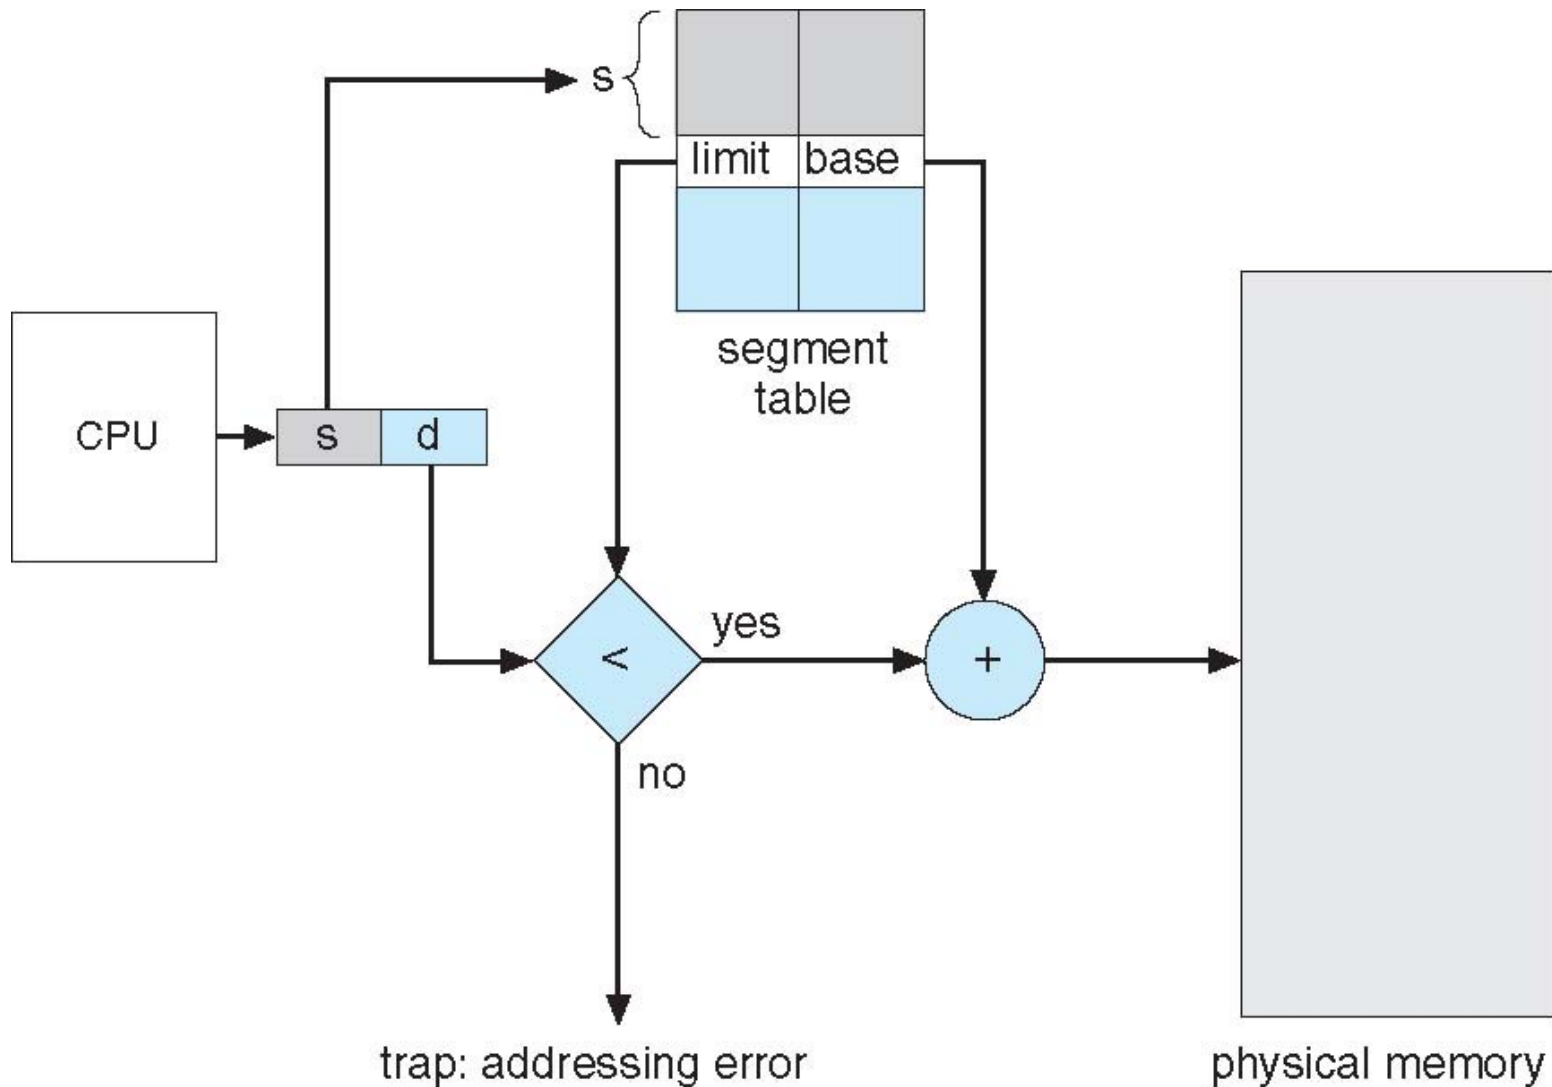
■ Segments Vary in Length

- \rightarrow Memory Allocation is a Dynamic Storage-Allocation Problem






Segmentation Hardware






Example of Segmentation



	limit	base
0	1000	1400
1	400	6300
2	400	4300
3	1100	3200
4	1000	4700

segment table





Segmentation vs. Variable-Sized Contiguous Allocation

■ Variable-Sized Contiguous Allocation

- Needs to bring entire process into memory
 - ▶ Both code, data, stack, ...

■ Segmentation

- Break program into different segments
- Brings segments to memory on demand
- No need to bring unused library methods or unused code & data segments to memory





Problems with Segmentation?

■ Example

- Code segment: 10MB
- Data segment: 100MB
- Stack segment: 20MB
- Lib segment: 10MB

■ Problem?

- Still significant amount of fragmentation
 - ▶ Both **external** & **internal** fragmentation





Paging

- Physical Address Space of a Process can be Non-Contiguous
 - Process is allocated physical memory whenever latter is available
 - Avoids external fragmentation
 - Avoids problem of varying sized memory chunks
- Divide Physical Memory into Fixed-sized blocks called **Frames**
 - Size is power of 2, between 512 bytes and 16 Mbytes





Paging (cont.)

- Divide Logical Memory into Blocks of Same Size called **Pages**
- Keep Track of all Free Frames
- To Run a Program of size **N** pages, Need to Find **N** Free Frames and Load Program
- Set up a **Page Table** to Translate Logical to Physical Addresses
- Backing Store Likewise Split into Pages
- Still have Internal Fragmentation






Address Translation Scheme

■ Address Generated by CPU divided into:

- **Page number** (p) – used as an index into a **page table** which contains base address of each page in physical memory
- **Page offset** (d) – combined with base address to define physical memory address that is sent to memory unit

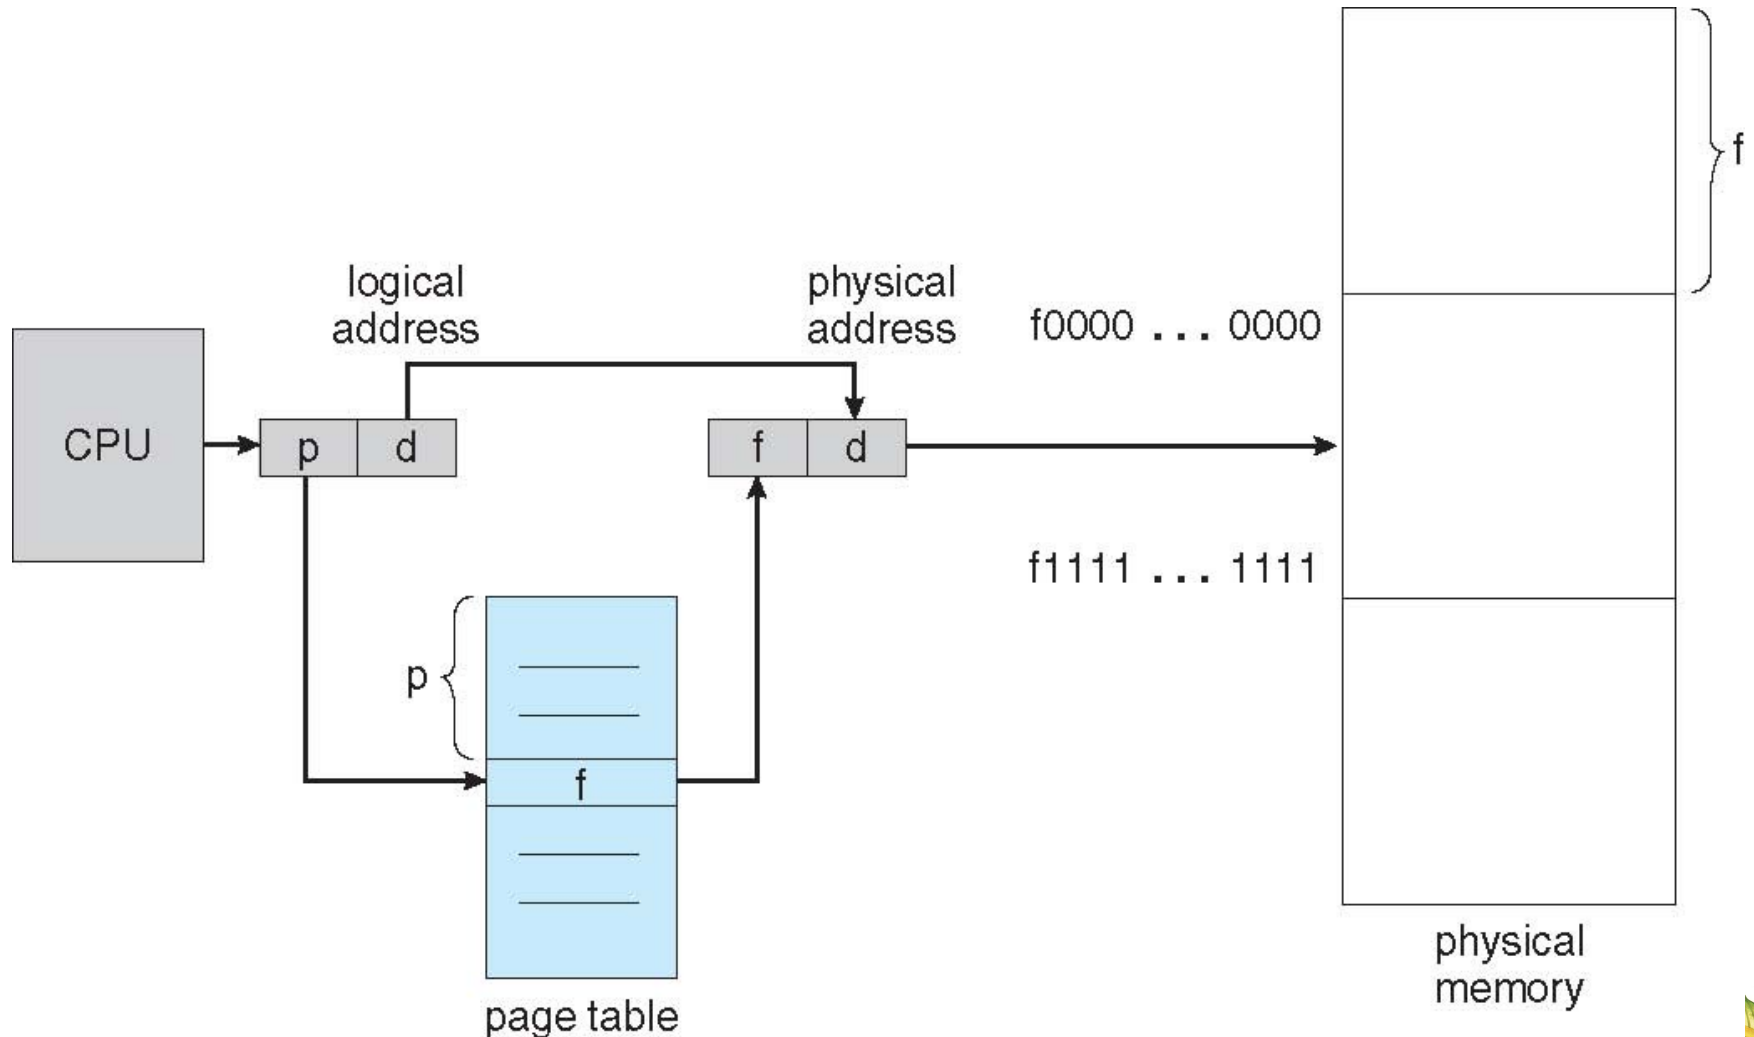


- ▶ Logical address space: 2^m
- ▶ Page size: 2^n






Paging Hardware





Paging Model of Logical and Physical Memory




logical
memory

0	1
1	4
2	3
3	7

page table

frame
number



physical
memory





Paging Example

0	a
1	b
2	c
3	d
4	e
5	f
6	g
7	h
8	i
9	j
10	k
11	l
12	m
13	n
14	o
15	p

logical memory

0	5
1	6
2	1
3	2

page table

0	
4	i j k l
8	m n o p
12	
16	
20	a b c d
24	e f g h
28	

physical memory

$n=2$ and $m=4$ 32-byte memory and 4-byte pages





Paging (cont.)

■ Calculating Internal Fragmentation

- Page size = 2,048 bytes
- Process size = 72,766 bytes
- 35 pages + 1,086 bytes
- Internal fragmentation of $2,048 - 1,086 = 962\text{B}$
- Worst case fragmentation = 1 frame – 1 byte
- On average fragmentation = $1 / 2$ frame size
- So small frame sizes desirable?
- But each page table entry takes memory to track
- Page sizes growing over time
 - ▶ Solaris supports two page sizes – 8 KB and 4 MB





Paging (cont.)

- Process View and Physical Memory now Very Different
- By Implementation, Process can only Access its Own Memory







Free Frames

free-frame list

14
13
18
20
15




(a)

Before allocation


free-frame list

15



0	14
1	13
2	18
3	20

new-process page table



(b)

After allocation





Implementation of Page Table

- Page Table Kept in Main Memory
- **Page-Table Base Register (PTBR)** Points to Page Table
- **Page-Table Length Register (PTLR)** Indicates Size of Page Table
- In this Scheme every Data/Instruction Access Requires **two Memory Accesses**
 - One for page table and one for data/instruction
- **Two Memory Access Problem** can be Solved by Use of a Special Fast-Lookup HW cache
 - Called **Associative Memory** or **Translation Look-Aside Buffers (TLBs)**





Implementation of Page Table (cont.)

■ Fully Associative TLB

- Fast 😊
- Very costly ☹️
- Typically 64~1024 entries

■ TLBs work similar to Caches

- Consists of two parts: **Key** (Tag) and a **Value**

■ TLB Functionality

- An incoming item compared with all keys simul.
- If found, corresponding value filed is returned
- If not found → **TLB miss**





Implementation of Page Table (cont.)

- Some TLBs Store **Address-Space Identifiers (ASIDs)** in each TLB entry
 - Uniquely identifies each process to provide address-space protection for that process
 - Otherwise need to flush at every context switch
- On a TLB Miss, Value Loaded into TLB for Faster Access Next Time
 - Replacement policies must be considered
 - Some entries **wired down** for permanent fast access (e.g., kernel code pages)





Associative Memory

■ Associative Memory – Parallel Search

Page #	Frame #

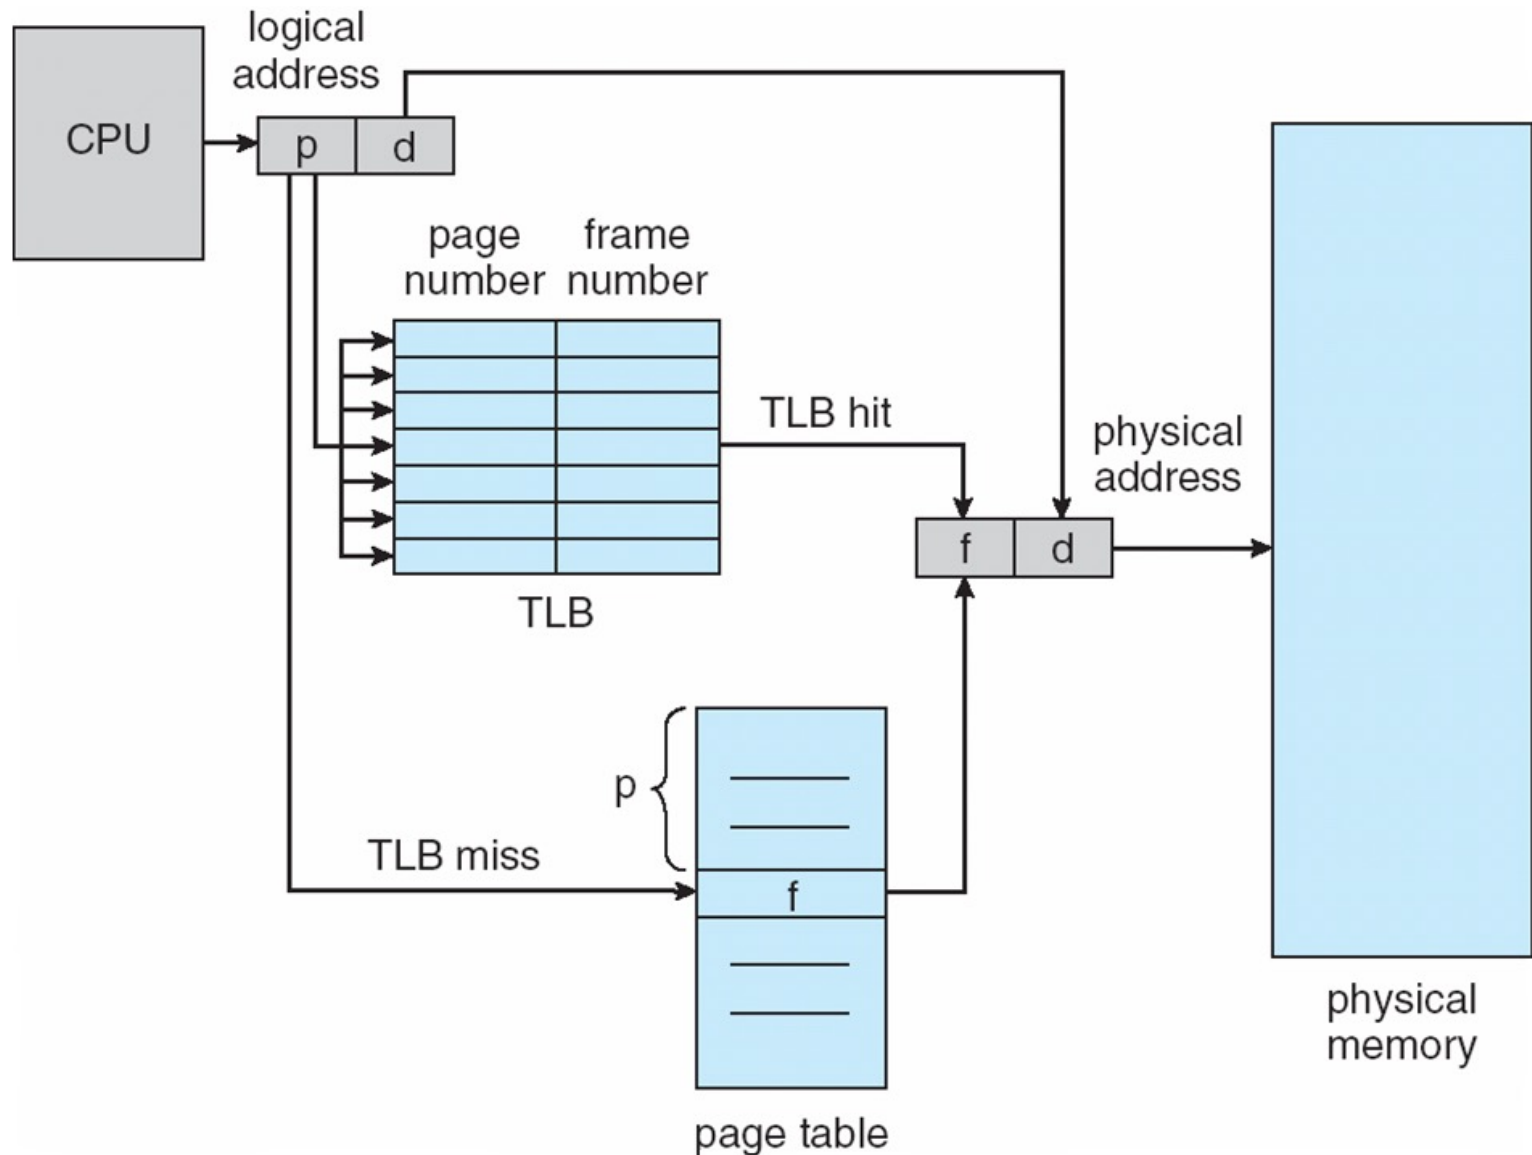
■ Address Translation (p, d)

- If p is in associative register, get **frame #** out
- Otherwise get frame # from page table in memory





Paging Hardware With TLB





Effective Access Time

■ Associative Lookup time unit (ε)

- Can be $< 10\%$ of memory access time

■ Hit ratio = α

- Hit ratio – percentage of times that a page number is found in associative registers; ratio related to number of associative registers

■ Consider $\alpha = 80\%$, $\varepsilon = 20\text{ns}$ for TLB search, 100ns for memory access

■ Effective Access Time (EAT)

$$\begin{aligned} \text{EAT} &= (1 + \varepsilon) \alpha + (2 + \varepsilon)(1 - \alpha) \\ &= 2 + \varepsilon - \alpha \end{aligned}$$





Effective Access Time (cont.)

■ Example 1:

- Consider $\alpha = 80\%$, $\varepsilon = 20\text{ns}$ for TLB search, 100ns for memory access
- $\text{EAT} = 0.80 \times 100 + 0.20 \times 200 = 120\text{ns}$

■ Example 2:

- Consider more realistic hit ratio
- $\alpha = 99\%$, $\varepsilon = 20\text{ns}$ for TLB search, 100ns for memory access
- $\text{EAT} = 0.99 \times 100 + 0.01 \times 200 = 101\text{ns}$





Memory Protection

- Memory Protection Implemented by Associating Protection bit with each Frame to Indicate if **Read-Only** or **Read-Write** Access is Allowed
 - Can also add more bits to indicate page **execute-only**
- **Valid-Invalid** bit Attached to each Entry in Page Table
 - “valid” indicates that associated page is in process’ logical address space, and is thus a legal page
 - “invalid” indicates that page is not in process’ logical address space
 - **Or** use **page-table length register (PTLR)**
- Any Violations Result in a **Trap** to Kernel





Valid (v) or Invalid (i) Bit In A Page Table

Example: a process with logical addresses ranging from 0 to 10468

00000	page 0
	page 1
	page 2
	page 3
	page 4
10,468	page 5
12,287	

frame number		valid-invalid bit
0	2	v
1	3	v
2	4	v
3	7	v
4	8	v
5	9	v
6	0	i
7	0	i

page table

0	
1	
2	page 0
3	page 1
4	page 2
5	
6	
7	page 3
8	page 4
9	page 5
	⋮
	page n

Any problem with this scheme?





Shared Pages

■ Shared code

- One copy of read-only (**reentrant**) code shared among processes (i.e., text editors, compilers, window systems)
- Similar to multiple threads sharing same process space
- Also useful for inter-process communication if sharing of read-write pages is allowed

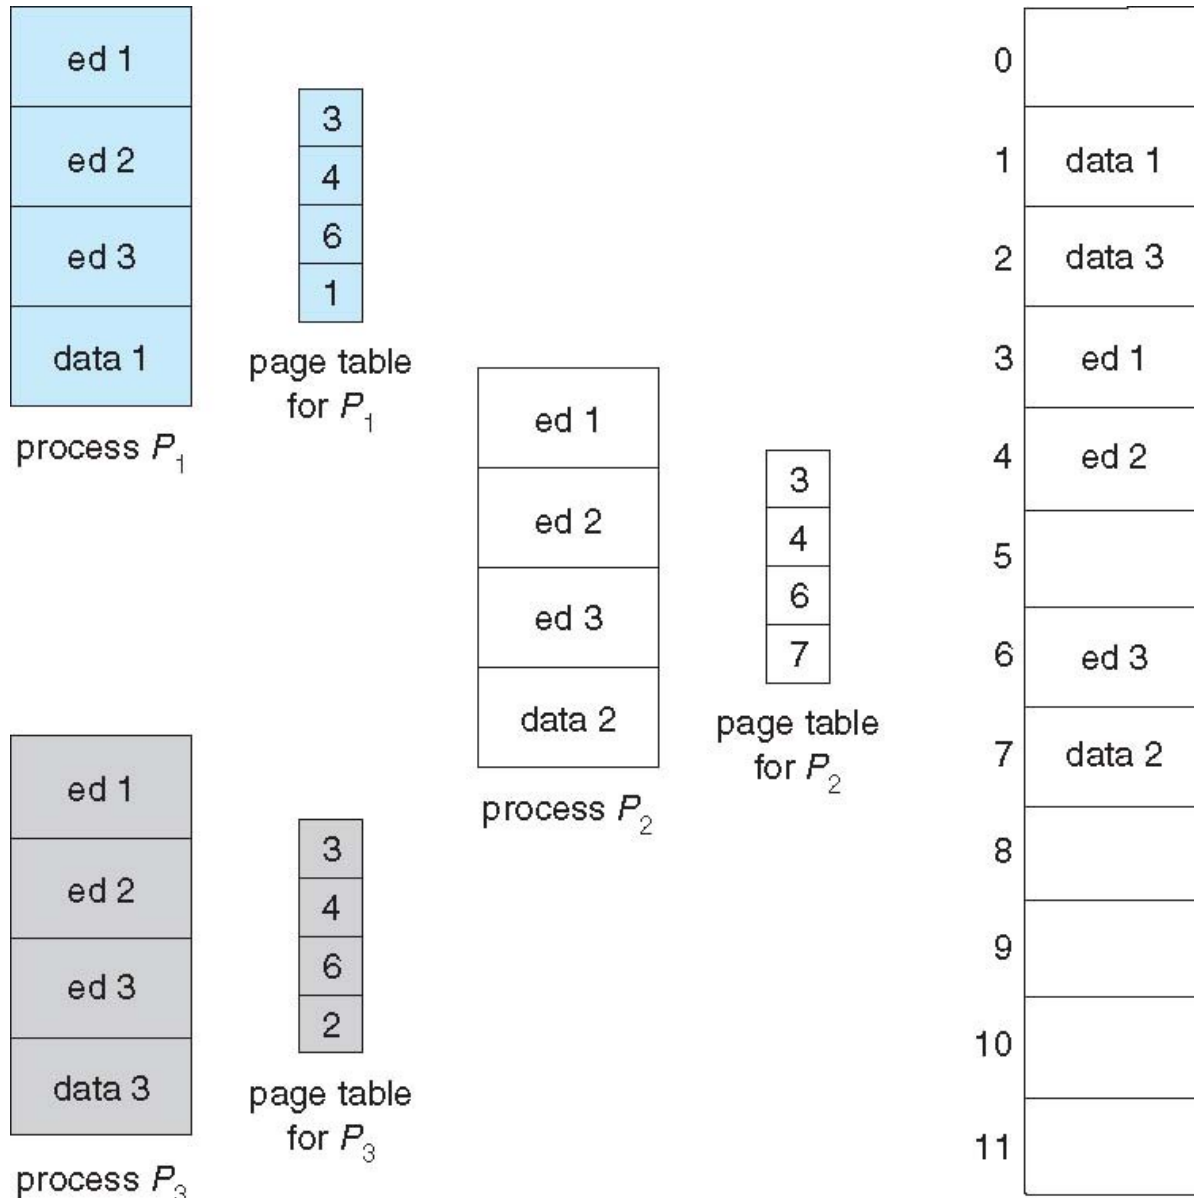
■ Private Code and Data

- Each process keeps a separate copy of code and data
- Pages for private code and data can appear anywhere in logical address space





Shared Pages Example





Structure of Page Table

■ Memory Structures for Paging can Get Huge using Straight-Forward Methods

- Consider a 32-bit logical address space as on modern computers
- Page size of 4 KB (2^{12})
- Page table would have 1 million entries ($2^{32} / 2^{12}$)
- If each entry is 4 bytes -> 4 MB of physical address space / memory for page table alone
 - ▶ That amount of memory used to cost a lot
 - ▶ Don't want to allocate that contiguously in main memory

■ Hierarchical Paging

■ Hashed Page Tables

■ Inverted Page Tables





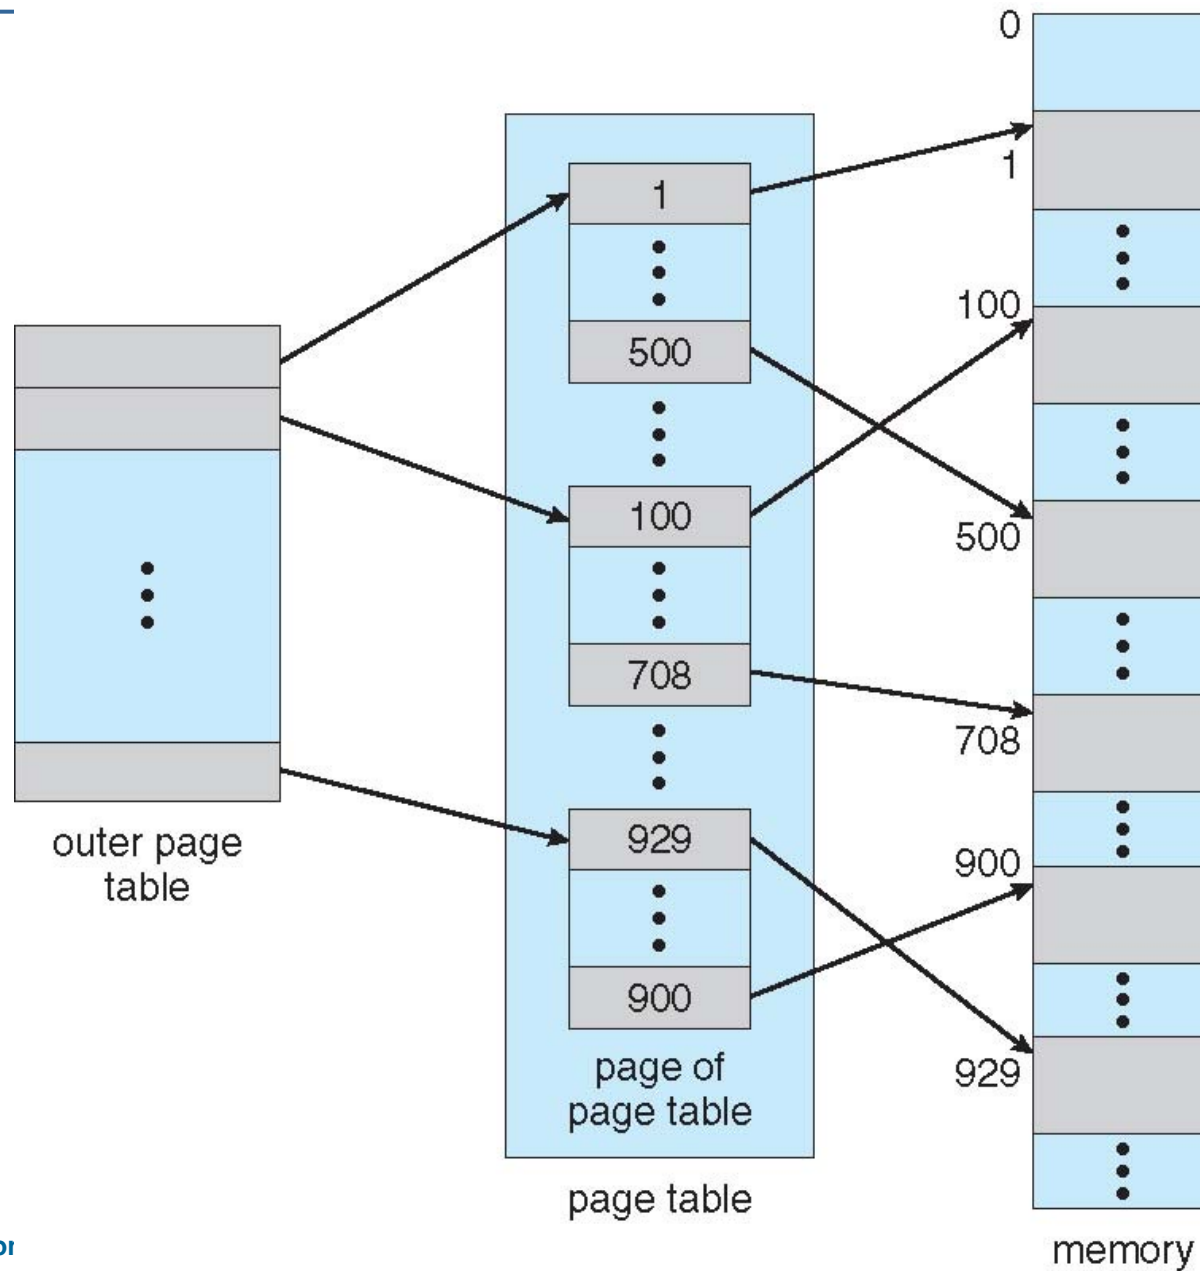
Hierarchical Page Tables

- Break up Logical Address Space into Multiple Page Tables
- A Simple Technique is a **Two-Level** Page table
- We then Page the page table





Two-Level Page-Table Scheme





Two-Level Paging Example

- A Logical Address (on 32-bit machine with 1K page size) Divided into:

- A page number consisting of 22 bits
- A page offset consisting of 10 bits

- Since Page Table is paged, page Number is further Divided into:

- A 12-bit page number
- A 10-bit page offset

page number		page offset
p_1	p_2	d
12	10	10

- Thus, a Logical Address is as follows:

- p_1 is an index into outer page table
- p_2 is displacement within page of inner page table


▶ Known as **forward-mapped page table**






Address-Translation Scheme

logical address




p_1 {




outer page table

p_2 {



page of page table

d {





64-bit Logical Address Space

■ Even two-level paging scheme not Sufficient

■ If Page Size is 4 KB (2^{12})



- Then page table has 2^{52} entries
- If 2-level scheme, inner page tables have 2^{10} 4-byte entries
- Address would look like
- Outer page table has 2^{42} entries or 2^{44} bytes
- One solution is to add a 2nd outer page table
- But in following example the 2nd outer page table is still 2^{34} bytes in size
 - ▶ Possibly 4 memory access to get to one physical memory location
→ hierarchical paging not appropriate for 64-bit addressing

outer page	inner page	page offset
p_1	p_2	d
42	10	12





Three-Level Paging Scheme






Hashed Page Tables

- Common in Address Spaces greater than 32 bits
- Virtual Page Number Hashed into a Page Table
 - This page table contains a chain of elements hashing to same location
- Each element contains (1) **virtual page number** (2) value of mapped page frame (3) a pointer to next element
- Virtual page numbers are compared in this chain searching for a match
 - If a match is found, corresponding physical frame is extracted





Hashed Page Table





Inverted Page Table

- **One Entry** for each Real Page of Memory
 - Rather than each process having a page table and keeping track of all possible logical pages, track all physical pages
- **Entry** consists of **virtual address** of page stored in that real memory location, with information about **process that owns** that page
- **Pros**
 - **Decreases memory needed** to store each page table





Inverted Page Table (cont.)

■ Cons

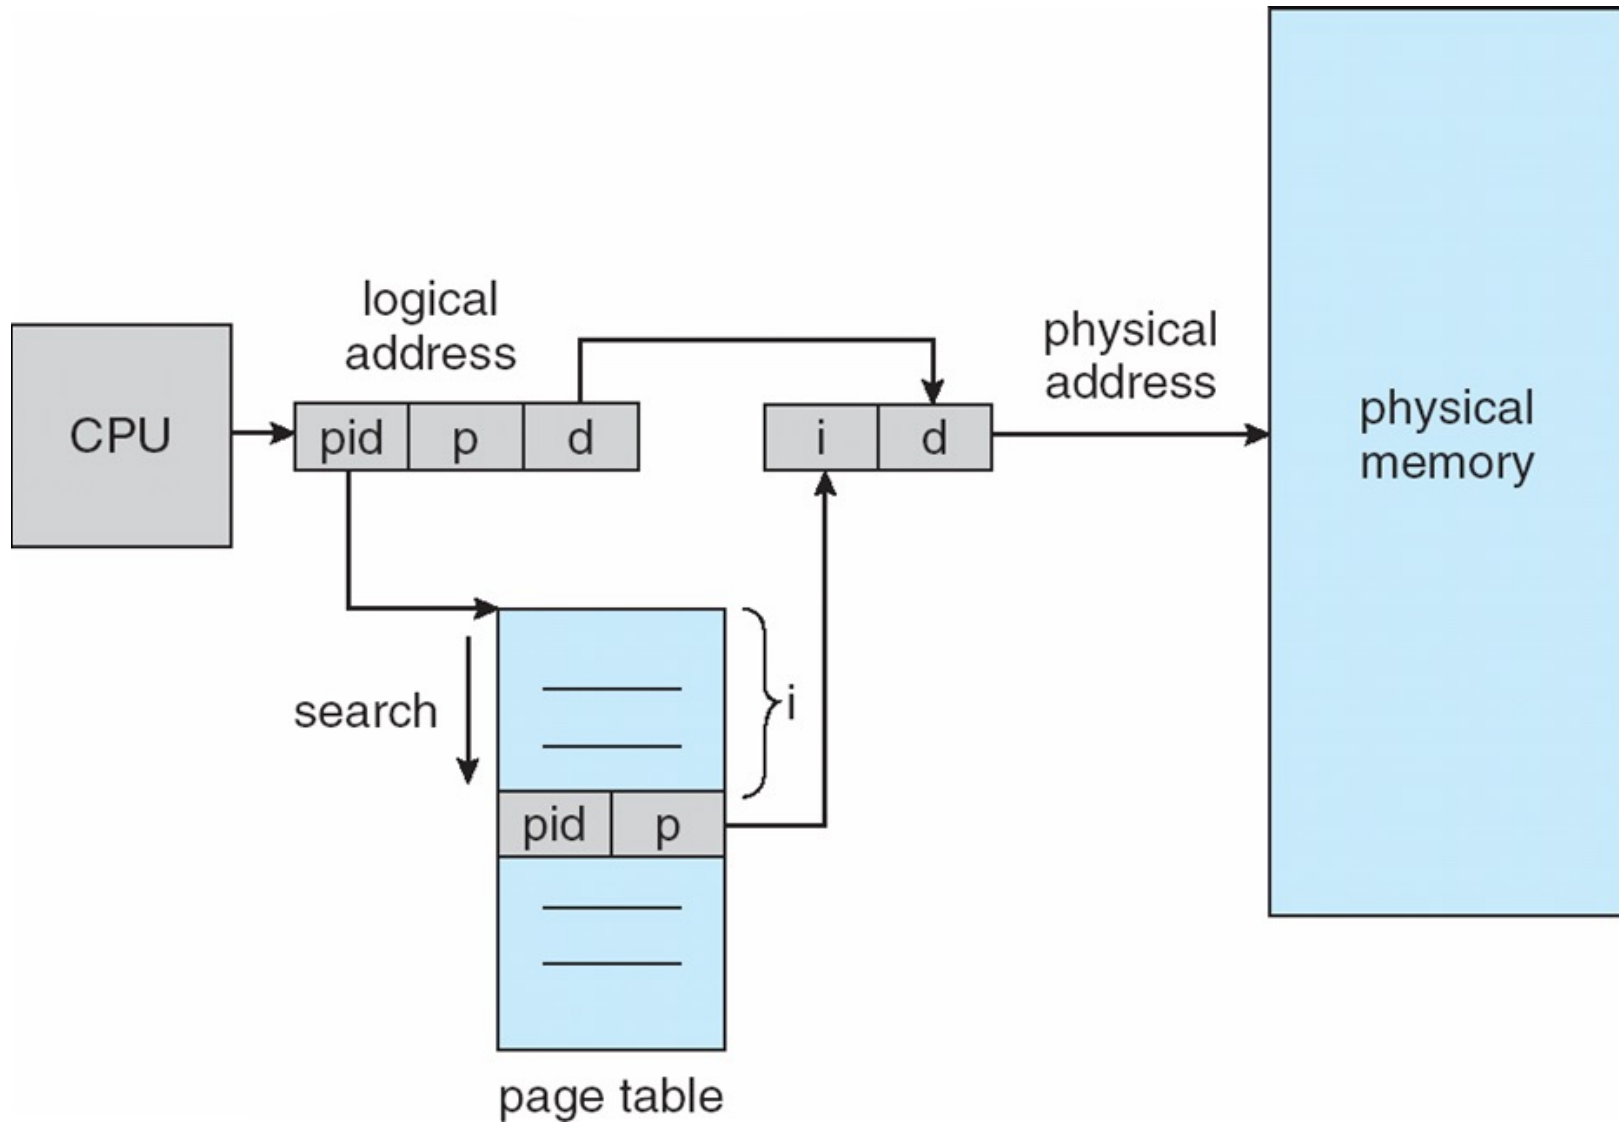
- Increases time needed to search table when a page reference occurs
 - ▶ Stored by physical addresses

- Use **Hash Table** to Limit Search to one — or at most a few — page-table entries
 - TLB can accelerate access





Inverted Page Table Architecture





Example: Intel 32 and 64-bit Architectures

- Dominant Industry Chips
- Pentium CPUs are 32-bit
 - Called IA-32 architecture
- Current Intel CPUs are 64-bit
 - Called IA-64 architecture





Example: Intel IA-32 Architecture

- Supports **both** segmentation and segmentation with paging
 - Each segment can be 4GB
 - Up to **16K** segments per process
 - Divided into two partitions
 - ▶ First partition of up to 8 K segments are **private** to process (kept in **local descriptor table (LDT)**)
 - ▶ Second partition of up to 8K segments **shared** among all processes (kept in **global descriptor table (GDT)**)






Example: Intel IA-32 Architecture (cont.)

■ CPU Generates Logical Address

- Selector given to segmentation unit
 - ▶ Which produces linear addresses

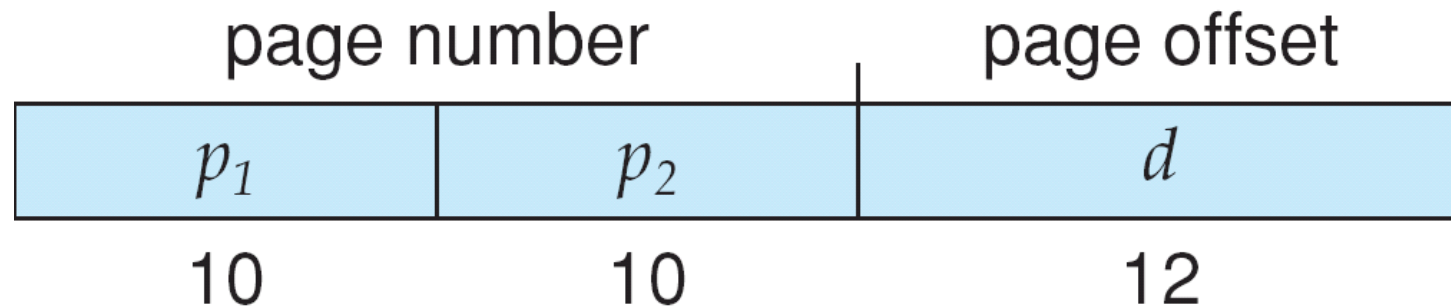



- Linear address given to paging unit
 - ▶ Which generates physical address in main memory
 - ▶ Paging units form equivalent of MMU
 - ▶ Pages sizes can be 4KB or 4MB





Logical to Physical Address Translation in IA-32





Reading Assignment

- Clustered Page Table
- Paging in Oracle SPARC Solaris
- Detailed Paging in IA-32 and IA-64





Hashed Page Tables (cont.)

- Variation for 64-bit Addresses is **Clustered Page Tables**
 - Similar to hashed but each entry refers to several pages (such as 16) rather than 1
 - Especially useful for **sparse** address spaces (where memory references are non-contiguous and scattered)





Oracle SPARC Solaris

- Consider modern, 64-bit OS example with tightly integrated HW
 - Goals are efficiency, low overhead
- Based on hashing, but more complex
- Two Hash Tables
 - One kernel and one for all user processes
 - Each maps memory addresses from virtual to physical memory
 - Each entry represents a contiguous area of mapped virtual memory
 - ▶ More efficient than having a separate hash-table entry for each page
 - Each entry has base address and span (indicating the number of pages the entry represents)





Oracle SPARC Solaris (cont.)

■ TLB holds translation table entries (TTEs) for fast hardware lookups

- A cache of TTEs reside in a translation storage buffer (TSB)

- ▶ Includes an entry per recently accessed page


■ Virtual address reference causes TLB search

- If miss, hardware walks the in-memory TSB looking for the TTE corresponding to the address

- ▶ If match found, the CPU copies the TSB entry into the TLB and translation completes

- ▶ If no match found, kernel interrupted to search the hash table

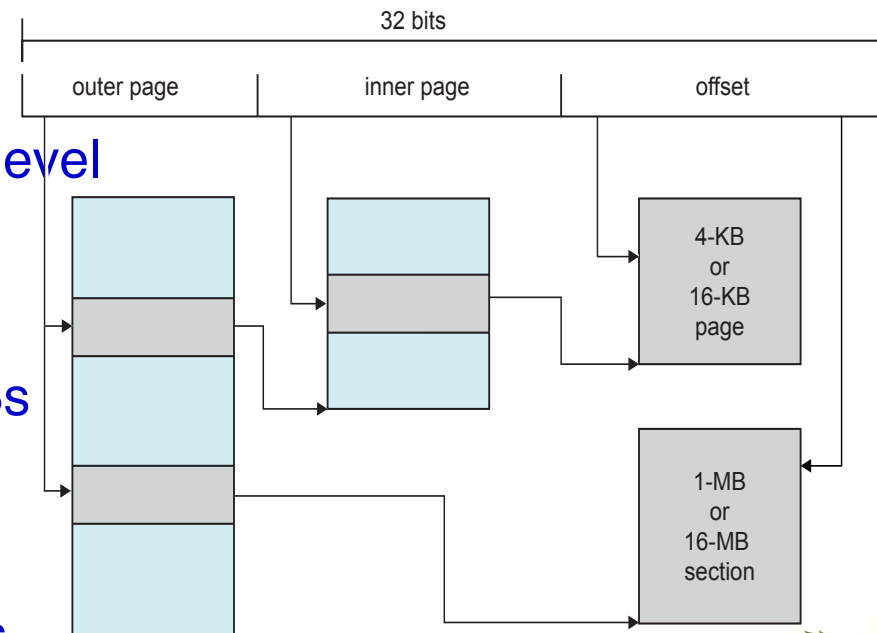
- The kernel then creates a TTE from the appropriate hash table and stores it in the TSB, Interrupt handler returns control to the MMU, which completes the address translation.





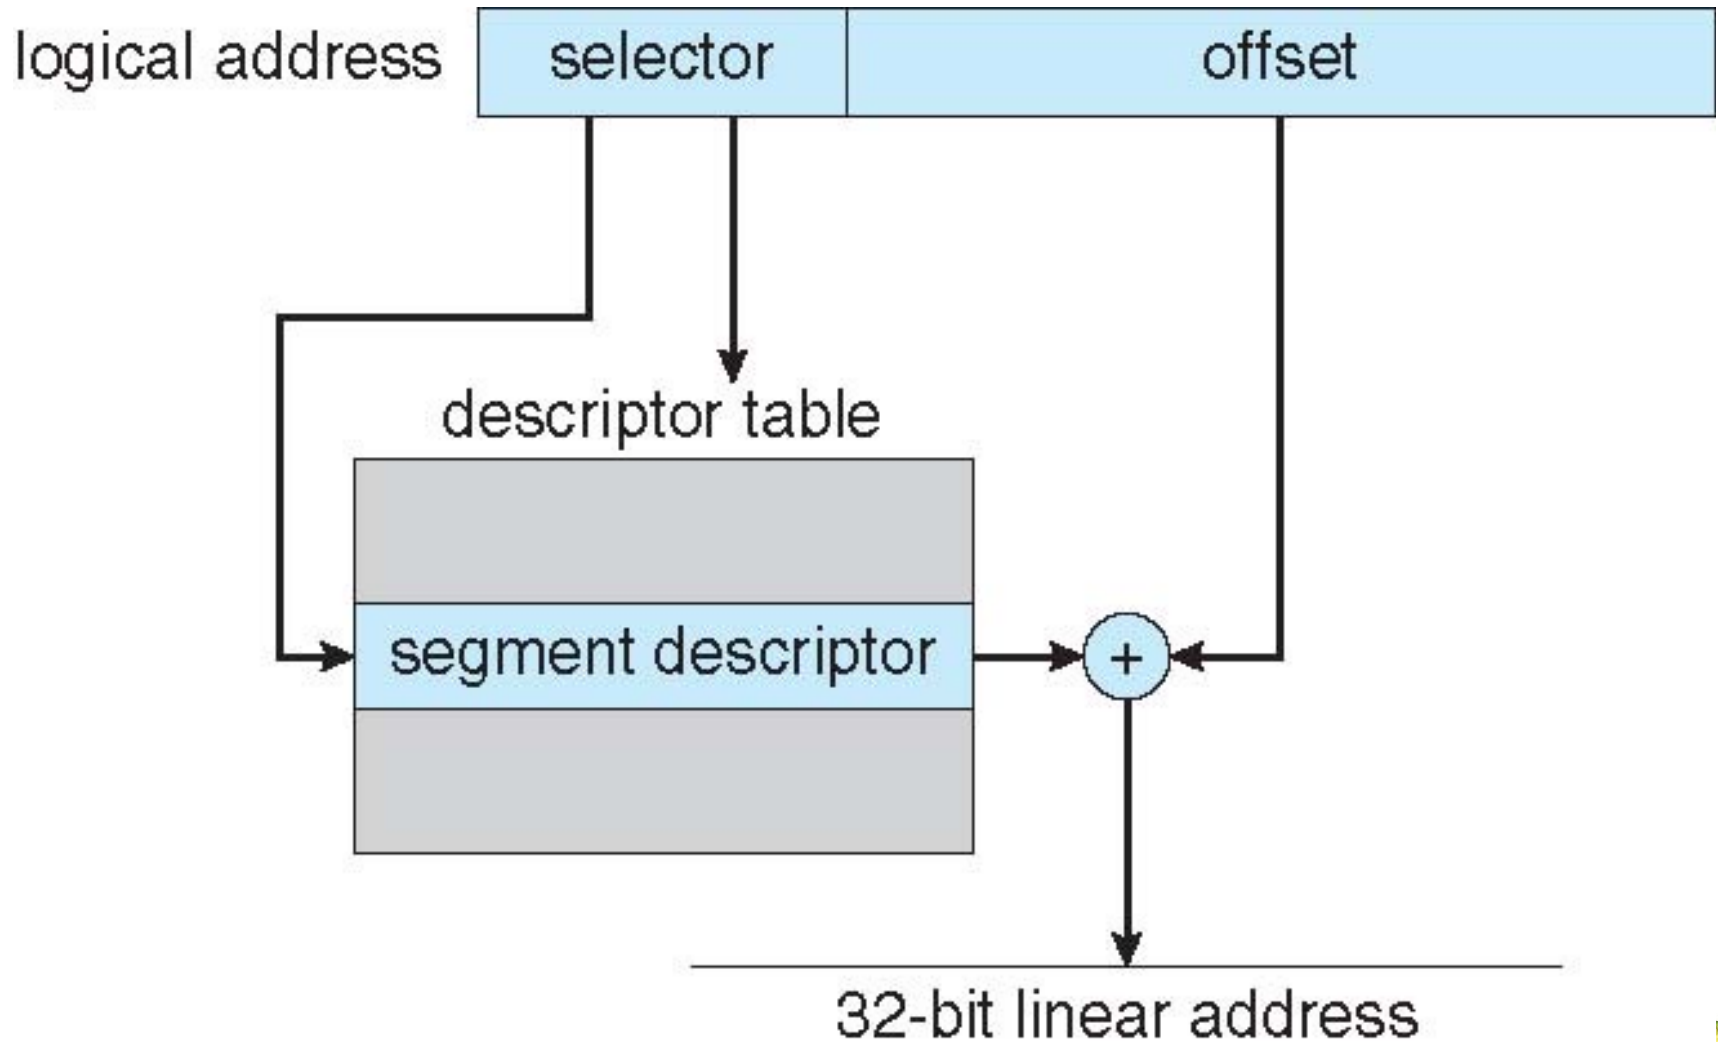
Example: ARM Architecture

- Dominant mobile platform chip (Apple iOS and Google Android devices for example)
- Modern, energy efficient, 32-bit CPU
- 4 KB and 16 KB pages
- 1 MB and 16 MB pages (termed **sections**)
- One-level paging for sections, two-level for smaller pages
- Two levels of TLBs
 - Outer level has two micro TLBs (one data, one instruction)
 - Inner is single main TLB
 - First inner is checked, on miss outers are checked, and on miss page table walk performed by CPU







Intel IA-32 Segmentation






Intel IA-32 Paging Architecture





Intel IA-32 Page Address Extensions



- 32-bit address limits led Intel to create **page address extension (PAE)**, allowing 32-bit apps access to more than 4GB of memory space
 - Paging went to a 3-level scheme
 - Top two bits refer to a **page directory pointer table**
 - Page-directory and page-table entries moved to 64-bits in size
 - Net effect is increasing address space to 36 bits – 64GB of physical memory





Intel x86-64

- Current generation Intel x86 architecture
- 64 bits is ginormous (> 16 exabytes)
- In practice only implement 48 bit addressing
 - Page sizes of 4KB, 2MB, 1GB
 - Four levels of paging hierarchy
- Can also use PAE so virtual addresses are 48 bits and physical addresses are 52 bits



End of Lecture 8

